

machine_learning_project

December 9, 2025

1 MOD10: Machine learning

Fall 2025 By : Arthur BOUTELIER, Alexandre BRENSKI, Paul LESAGE, Adrien RIVET Instructor: Mohammed A. Shehab

GITHUB : <https://github.com/Arthur-Boutelier/machineLearningCybersecurity>

2 Introduction

This project implements a complete machine learning pipeline for network intrusion detection using a dataset downloaded from Kaggle. The pipeline covers: - Data acquisition - Cleaning & preprocessing - Feature analysis - Handling categorical & numerical attributes - Scaling & encoding - Train/test splitting - Model training - Evaluation (confusion matrix, classification report, heatmap visualization)

2.1 Environment Setup & Importing Dependencies

```
[ ]: ! pip install -r requirements.txt

Requirement already satisfied: alembic==1.17.2 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 1)) (1.17.2)
Requirement already satisfied: annotated-doc==0.0.4 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 2)) (0.0.4)
Requirement already satisfied: annotated-types==0.7.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 3)) (0.7.0)
Requirement already satisfied: anyio==4.12.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 4)) (4.12.0)
Requirement already satisfied: asttokens==3.0.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 5)) (3.0.0)
Requirement already satisfied: blinker==1.9.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 6)) (1.9.0)
Requirement already satisfied: cachetools==6.2.2 in c:\users\arthu\onedrive\efre
```

```
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 7)) (6.2.2)
Requirement already satisfied: certifi==2025.11.12 in c:\users\arthu\onedrive\ef
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from
-r requirements.txt (line 8)) (2025.11.12)
Requirement already satisfied: cffi==2.0.0 in c:\users\arthu\onedrive\efrei\i1\s
5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 9)) (2.0.0)
Requirement already satisfied: charset-normalizer==3.4.4 in c:\users\arthu\onedr
ive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages
(from -r requirements.txt (line 10)) (3.4.4)
Requirement already satisfied: click==8.3.1 in c:\users\arthu\onedrive\efrei\i1\
s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 11)) (8.3.1)
Requirement already satisfied: cloudpickle==3.1.2 in c:\users\arthu\onedrive\efr
ei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 12)) (3.1.2)
Requirement already satisfied: colorama==0.4.6 in c:\users\arthu\onedrive\efrei\
i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 13)) (0.4.6)
Requirement already satisfied: comm==0.2.3 in c:\users\arthu\onedrive\efrei\i1\s
5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 14)) (0.2.3)
Requirement already satisfied: contourpy==1.3.3 in c:\users\arthu\onedrive\efrei
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 15)) (1.3.3)
Requirement already satisfied: cryptography==46.0.3 in c:\users\arthu\onedrive\ef
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 16)) (46.0.3)
Requirement already satisfied: cycler==0.12.1 in c:\users\arthu\onedrive\efrei\i
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 17)) (0.12.1)
Requirement already satisfied: databricks-sdk==0.73.0 in c:\users\arthu\onedrive
\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 18)) (0.73.0)
Requirement already satisfied: debugpy==1.8.17 in c:\users\arthu\onedrive\efrei\
i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 19)) (1.8.17)
Requirement already satisfied: decorator==5.2.1 in c:\users\arthu\onedrive\efrei
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 20)) (5.2.1)
Requirement already satisfied: docker==7.1.0 in c:\users\arthu\onedrive\efrei\i1
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 21)) (7.1.0)
Requirement already satisfied: executing==2.2.1 in c:\users\arthu\onedrive\efrei
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 22)) (2.2.1)
Requirement already satisfied: fastapi==0.124.0 in c:\users\arthu\onedrive\efrei
```

```
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 23)) (0.124.0)
Requirement already satisfied: Flask==3.1.2 in c:\users\arthu\onedrive\efrei\i1\
s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 24)) (3.1.2)
Requirement already satisfied: flask-cors==6.0.1 in c:\users\arthu\onedrive\efre
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 25)) (6.0.1)
Requirement already satisfied: fonttools==4.60.1 in c:\users\arthu\onedrive\efre
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 26)) (4.60.1)
Requirement already satisfied: gitdb==4.0.12 in c:\users\arthu\onedrive\efrei\i1
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 27)) (4.0.12)
Requirement already satisfied: GitPython==3.1.45 in c:\users\arthu\onedrive\efre
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 28)) (3.1.45)
Requirement already satisfied: google-auth==2.43.0 in c:\users\arthu\onedrive\ef
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from
-r requirements.txt (line 29)) (2.43.0)
Requirement already satisfied: graphene==3.4.3 in c:\users\arthu\onedrive\efrei\
i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 30)) (3.4.3)
Requirement already satisfied: graphql-core==3.2.7 in c:\users\arthu\onedrive\ef
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from
-r requirements.txt (line 31)) (3.2.7)
Requirement already satisfied: graphql-relay==3.2.0 in c:\users\arthu\onedrive\ef
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from
-r requirements.txt (line 32)) (3.2.0)
Requirement already satisfied: greenlet==3.3.0 in c:\users\arthu\onedrive\efrei\
i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 33)) (3.3.0)
Requirement already satisfied: h11==0.16.0 in c:\users\arthu\onedrive\efrei\i1\s
5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 34)) (0.16.0)
Requirement already satisfied: huey==2.5.5 in c:\users\arthu\onedrive\efrei\i1\s
5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 35)) (2.5.5)
Requirement already satisfied: idna==3.11 in c:\users\arthu\onedrive\efrei\i1\s5
\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 36)) (3.11)
Requirement already satisfied: imbalanced-learn==0.14.0 in c:\users\arthu\onedri
ve\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages
(from -r requirements.txt (line 37)) (0.14.0)
Requirement already satisfied: imblearn==0.0 in c:\users\arthu\onedrive\efrei\i1
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 38)) (0.0)
Requirement already satisfied: importlib_metadata==8.7.0 in c:\users\arthu\onedr
```

```
ive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages  
(from -r requirements.txt (line 39)) (8.7.0)  
Requirement already satisfied: ipykernel==7.1.0 in c:\users\arthu\onedrive\efrei  
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 40)) (7.1.0)  
Requirement already satisfied: ipython==9.7.0 in c:\users\arthu\onedrive\efrei\i  
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 41)) (9.7.0)  
Requirement already satisfied: ipython_pygments_lexers==1.1.1 in c:\users\arthu\  
onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-  
packages (from -r requirements.txt (line 42)) (1.1.1)  
Requirement already satisfied: itsdangerous==2.2.0 in c:\users\arthu\onedrive\ef  
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from  
-r requirements.txt (line 43)) (2.2.0)  
Requirement already satisfied: jedi==0.19.2 in c:\users\arthu\onedrive\efrei\i1\  
s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 44)) (0.19.2)  
Requirement already satisfied: Jinja2==3.1.6 in c:\users\arthu\onedrive\efrei\i1\  
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 45)) (3.1.6)  
Requirement already satisfied: joblib==1.5.2 in c:\users\arthu\onedrive\efrei\i1\  
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 46)) (1.5.2)  
Requirement already satisfied: jupyter_client==8.6.3 in c:\users\arthu\onedrive\ef  
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from  
-r requirements.txt (line 47)) (8.6.3)  
Requirement already satisfied: jupyter_core==5.9.1 in c:\users\arthu\onedrive\ef  
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from  
-r requirements.txt (line 48)) (5.9.1)  
Requirement already satisfied: kagglehub==0.3.13 in c:\users\arthu\onedrive\efre  
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 49)) (0.3.13)  
Requirement already satisfied: kiwisolver==1.4.9 in c:\users\arthu\onedrive\efre  
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 50)) (1.4.9)  
Requirement already satisfied: Mako==1.3.10 in c:\users\arthu\onedrive\efrei\i1\  
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 51)) (1.3.10)  
Requirement already satisfied: MarkupSafe==3.0.3 in c:\users\arthu\onedrive\efre  
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 52)) (3.0.3)  
Requirement already satisfied: matplotlib==3.10.7 in c:\users\arthu\onedrive\efr  
ei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 53)) (3.10.7)  
Requirement already satisfied: matplotlib-inline==0.2.1 in c:\users\arthu\onedri  
ve\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages  
(from -r requirements.txt (line 54)) (0.2.1)  
Requirement already satisfied: mlflow==3.6.0 in c:\users\arthu\onedrive\efrei\i1
```

```
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 55)) (3.6.0)
Requirement already satisfied: mlflow-skinny==3.6.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 56)) (3.6.0)
Requirement already satisfied: mlflow-tracing==3.6.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 57)) (3.6.0)
Requirement already satisfied: nest-asyncio==1.6.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 58)) (1.6.0)
Requirement already satisfied: numpy==2.3.4 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 59)) (2.3.4)
Requirement already satisfied: opentelemetry-api==1.39.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 60)) (1.39.0)
Requirement already satisfied: opentelemetry-proto==1.39.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 61)) (1.39.0)
Requirement already satisfied: opentelemetry-sdk==1.39.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 62)) (1.39.0)
Requirement already satisfied: opentelemetry-semantic-conventions==0.60b0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 63)) (0.60b0)
Requirement already satisfied: packaging==25.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 64)) (25.0)
Requirement already satisfied: pandas==2.3.3 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 65)) (2.3.3)
Requirement already satisfied: parso==0.8.5 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 66)) (0.8.5)
Requirement already satisfied: pillow==12.0.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 67)) (12.0.0)
Requirement already satisfied: platformdirs==4.5.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 68)) (4.5.0)
Requirement already satisfied: prompt_toolkit==3.0.52 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 69)) (3.0.52)
Requirement already satisfied: protobuf==6.33.2 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r requirements.txt (line 70)) (6.33.2)
Requirement already satisfied: psutil==7.1.3 in c:\users\arthu\onedrive\efrei\i1
```

```
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 71)) (7.1.3)
Requirement already satisfied: pure_eval==0.2.3 in c:\users\arthu\onedrive\efrei
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 72)) (0.2.3)
Requirement already satisfied: pyarrow==22.0.0 in c:\users\arthu\onedrive\efrei\
i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 73)) (22.0.0)
Requirement already satisfied: pyasn1==0.6.1 in c:\users\arthu\onedrive\efrei\i1
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 74)) (0.6.1)
Requirement already satisfied: pyasn1_modules==0.4.2 in c:\users\arthu\onedrive\
efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from
-r requirements.txt (line 75)) (0.4.2)
Requirement already satisfied: pycparser==2.23 in c:\users\arthu\onedrive\efrei\
i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 76)) (2.23)
Requirement already satisfied: pydantic==2.12.5 in c:\users\arthu\onedrive\efrei
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 77)) (2.12.5)
Requirement already satisfied: pydantic_core==2.41.5 in c:\users\arthu\onedrive\
efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from
-r requirements.txt (line 78)) (2.41.5)
Requirement already satisfied: Pygments==2.19.2 in c:\users\arthu\onedrive\efrei
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 79)) (2.19.2)
Requirement already satisfied: pyparsing==3.2.5 in c:\users\arthu\onedrive\efrei
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 80)) (3.2.5)
Requirement already satisfied: python-dateutil==2.9.0.post0 in c:\users\arthu\on
edrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-
packages (from -r requirements.txt (line 81)) (2.9.0.post0)
Requirement already satisfied: python-dotenv==1.2.1 in c:\users\arthu\onedrive\ef
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from
-r requirements.txt (line 82)) (1.2.1)
Requirement already satisfied: pytz==2025.2 in c:\users\arthu\onedrive\efrei\i1\
s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 83)) (2025.2)
Requirement already satisfied: pywin32==311 in c:\users\arthu\onedrive\efrei\i1\
s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 84)) (311)
Requirement already satisfied: PyYAML==6.0.3 in c:\users\arthu\onedrive\efrei\i1
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 85)) (6.0.3)
Requirement already satisfied: pyzmq==27.1.0 in c:\users\arthu\onedrive\efrei\i1
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 86)) (27.1.0)
Requirement already satisfied: requests==2.32.5 in c:\users\arthu\onedrive\efrei
```

```
\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 87)) (2.32.5)
Requirement already satisfied: rsa==4.9.1 in c:\users\arthu\onedrive\efrei\i1\s5
\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 88)) (4.9.1)
Requirement already satisfied: scikit-learn==1.7.2 in c:\users\arthu\onedrive\ef
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from
-r requirements.txt (line 89)) (1.7.2)
Requirement already satisfied: scipy==1.16.3 in c:\users\arthu\onedrive\efrei\i1
\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 90)) (1.16.3)
Requirement already satisfied: seaborn==0.13.2 in c:\users\arthu\onedrive\efrei\
i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 91)) (0.13.2)
Requirement already satisfied: six==1.17.0 in c:\users\arthu\onedrive\efrei\i1\s
5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 92)) (1.17.0)
Requirement already satisfied: smmap==5.0.2 in c:\users\arthu\onedrive\efrei\i1\
s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 93)) (5.0.2)
Requirement already satisfied: SQLAlchemy==2.0.44 in c:\users\arthu\onedrive\efr
ei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 94)) (2.0.44)
Requirement already satisfied: sqlparse==0.5.4 in c:\users\arthu\onedrive\efrei\
i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 95)) (0.5.4)
Requirement already satisfied: stack-data==0.6.3 in c:\users\arthu\onedrive\efre
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 96)) (0.6.3)
Requirement already satisfied: starlette==0.50.0 in c:\users\arthu\onedrive\efre
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 97)) (0.50.0)
Requirement already satisfied: threadpoolctl==3.6.0 in c:\users\arthu\onedrive\ef
rei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 98)) (3.6.0)
Requirement already satisfied: tornado==6.5.2 in c:\users\arthu\onedrive\efrei\i
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 99)) (6.5.2)
Requirement already satisfied: tqdm==4.67.1 in c:\users\arthu\onedrive\efrei\i1\
s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 100)) (4.67.1)
Requirement already satisfied: traitlets==5.14.3 in c:\users\arthu\onedrive\efre
i\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r
requirements.txt (line 101)) (5.14.3)
Requirement already satisfied: typing-inspection==0.4.2 in c:\users\arthu\onedri
ve\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages
(from -r requirements.txt (line 102)) (0.4.2)
Requirement already satisfied: typing_extensions==4.15.0 in c:\users\arthu\onedr
```

```
ive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages  
(from -r requirements.txt (line 103)) (4.15.0)  
Requirement already satisfied: tzdata==2025.2 in c:\users\arthu\onedrive\efrei\i  
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 104)) (2025.2)  
Requirement already satisfied: urllib3==2.5.0 in c:\users\arthu\onedrive\efrei\i  
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 105)) (2.5.0)  
Requirement already satisfied: uvicorn==0.38.0 in c:\users\arthu\onedrive\efrei\i  
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 106)) (0.38.0)  
Requirement already satisfied: waitress==3.0.2 in c:\users\arthu\onedrive\efrei\i  
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 107)) (3.0.2)  
Requirement already satisfied: wcwidth==0.2.14 in c:\users\arthu\onedrive\efrei\i  
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 108)) (0.2.14)  
Requirement already satisfied: Werkzeug==3.1.4 in c:\users\arthu\onedrive\efrei\i  
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 109)) (3.1.4)  
Requirement already satisfied: xgboost==3.1.2 in c:\users\arthu\onedrive\efrei\i  
1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 110)) (3.1.2)  
Requirement already satisfied: zipp==3.23.0 in c:\users\arthu\onedrive\efrei\i1\s5\machinelearning\machinelearningcybersecurity\lib\site-packages (from -r  
requirements.txt (line 111)) (3.23.0)
```

[notice] A new release of pip is available: 25.2 -> 25.3

[notice] To update, run: python.exe -m pip install --upgrade pip

This block imports all required Python libraries used throughout the project.

1/ General-purpose libraries - os: Interacting with the operating system (listing files, building paths). - pandas: For data loading, manipulation, DataFrame operations. - numpy: For numerical computation, array processing. - seaborn, matplotlib.pyplot: For plotting charts and visual analysis.

2/ Machine Learning - StandardScaler: Normalizes numerical features (important for models like Logistic Regression). - LabelEncoder: Converts categorical labels into integers. - train_test_split: Splits data into training/testing sets. - LogisticRegression: A baseline ML classification model. - classification_report, confusion_matrix: Evaluation metrics. - DecisionTreeClassifier: A simple tree-based ML model.

3/ Advanced ML - xgboost: Gradient boosting model, high-performance. - mlflow: Used for experiment tracking and model versioning.

4/ Purpose - This central import block prepares all tools needed for: - Data cleaning - Preprocessing - Machine learning - Visualization - Model evaluation - Experiment logging

```
[42]: import os
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, StandardScaler
import kagglehub
from sklearn.model_selection import train_test_split, GridSearchCV
from imblearn.over_sampling import SMOTE
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix,
    accuracy_score, f1_score
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, VotingClassifier,
    StackingClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
import mlflow as mlf
import xgboost as xgb
```

3 1. Data Preprocessing and EDA

3.1 Load and clean dataset

Load and clean the dataset, handling missing values, normalizing or scaling numerical features, and encoding categorical variables.

This block downloads the network intrusion Dataset from Kaggle through kagglehub.

```
[48]: # Download latest version
path = kagglehub.dataset_download("chethuhn/network-intrusion-dataset")
print("Path to dataset files:", path)
models = []
```

Path to dataset files:
C:\Users\arthu\.cache\kagglehub\datasets\chethuhn\network-intrusion-dataset\versions\1

This block lists all files downloaded from Kaggle.

```
[49]: # see list of downloaded dataset files
list_of_files = os.listdir(path)
print("Files in dataset:", list_of_files)
```

Files in dataset: ['Friday-WorkingHours-Afternoon-DDos.pcap_ISCX.csv', 'Friday-WorkingHours-Afternoon-PortScan.pcap_ISCX.csv', 'Friday-WorkingHours-Morning.pcap_ISCX.csv', 'Monday-WorkingHours.pcap_ISCX.csv', 'Thursday-WorkingHours-Afternoon-Infiltration.pcap_ISCX.csv', 'Thursday-WorkingHours-Morning-WebAttacks.pcap_ISCX.csv', 'Tuesday-WorkingHours.pcap_ISCX.csv',

```
'Wednesday-workingHours.pcap_ISCX.csv']
```

Loads the dataset into a pandas DataFrame. df.head() shows the first 5 rows to verify correct loading.

```
[50]: # Load first dataset
```

```
df = pd.read_csv(f'{path}/{list_of_files[0]}')  
df.head()
```

```
[50]:    Destination Port  Flow Duration  Total Fwd Packets  \  
0           54865            3                 2  
1           55054          109                 1  
2           55055            52                 1  
3           46236            34                 1  
4           54863            3                 2  
  
   Total Backward Packets  Total Length of Fwd Packets  \  
0                  0                   12  
1                  1                   6  
2                  1                   6  
3                  1                   6  
4                  0                   12  
  
   Total Length of Bwd Packets  Fwd Packet Length Max  \  
0                  0                   6  
1                  6                   6  
2                  6                   6  
3                  6                   6  
4                  0                   6  
  
   Fwd Packet Length Min  Fwd Packet Length Mean  Fwd Packet Length Std  \  
0                  6                  6.0                0.0  
1                  6                  6.0                0.0  
2                  6                  6.0                0.0  
3                  6                  6.0                0.0  
4                  6                  6.0                0.0  
  
   ...  min_seg_size_forward  Active Mean  Active Std  Active Max  \  
0 ...              20        0.0        0.0        0  
1 ...              20        0.0        0.0        0  
2 ...              20        0.0        0.0        0  
3 ...              20        0.0        0.0        0  
4 ...              20        0.0        0.0        0  
  
   Active Min  Idle Mean  Idle Std  Idle Max  Idle Min  Label  
0          0      0.0      0.0        0        0  BENIGN  
1          0      0.0      0.0        0        0  BENIGN  
2          0      0.0      0.0        0        0  BENIGN
```

```

3          0      0.0      0.0      0      0  BENIGN
4          0      0.0      0.0      0      0  BENIGN

```

[5 rows x 79 columns]

This code loads every CSV file in the dataset, adds a column identifying the file of origin, stores each loaded file in a list, and finally concatenates all of them into one unified DataFrame while printing its final shape and previewing the first rows.

We concatenate all datasets because they have the same columns and represent different days of the same measurements.

We concatenate all datasets because they have the same columns and represent different days of the same measurements.

```
[51]: # concatenate all dataset files
dataframes = []
for i, file in enumerate(list_of_files):
    df_part = pd.read_csv(f"{path}/{file}")
    df_part['source_file'] = i # optional: add a column to identify source file
    dataframes.append(df_part)

# concatenate all dataframes into a single dataframe
df = pd.concat(dataframes, axis=0, ignore_index=True)
print("Combined dataset shape:", df.shape)
df.head()
```

Combined dataset shape: (2830743, 80)

```
[51]:   Destination Port  Flow Duration  Total Fwd Packets \
0           54865            3                 2
1           55054          109                 1
2           55055            52                 1
3           46236            34                 1
4           54863            3                 2

   Total Backward Packets  Total Length of Fwd Packets \
0                      0                     12
1                      1                     6
2                      1                     6
3                      1                     6
4                      0                     12

   Total Length of Bwd Packets  Fwd Packet Length Max \
0                      0                     6
1                      6                     6
2                      6                     6
3                      6                     6
4                      0                     6
```

	Fwd Packet Length Min	Fwd Packet Length Mean	Fwd Packet Length Std	\
0	6	6.0	0.0	
1	6	6.0	0.0	
2	6	6.0	0.0	
3	6	6.0	0.0	
4	6	6.0	0.0	

	... Active Mean	Active Std	Active Max	Active Min	Idle Mean	\
0 ...	0.0	0.0	0	0	0.0	
1 ...	0.0	0.0	0	0	0.0	
2 ...	0.0	0.0	0	0	0.0	
3 ...	0.0	0.0	0	0	0.0	
4 ...	0.0	0.0	0	0	0.0	

	Idle Std	Idle Max	Idle Min	Label	source_file	
0	0.0	0	0	BENIGN	0	
1	0.0	0	0	BENIGN	0	
2	0.0	0	0	BENIGN	0	
3	0.0	0	0	BENIGN	0	
4	0.0	0	0	BENIGN	0	

[5 rows x 80 columns]

This code enables a test mode that randomly selects a smaller 50,000-row subset of the full dataset to speed up development and experimentation, and then prints the reduced dataset's shape.

```
[52]: # define test mode, to work with a smaller subset during development, instead ↴ of 2.8M rows
TEST_MODE = True
if TEST_MODE:
    df = df.sample(n=50000, random_state=42)
    print("Test mode: using smaller dataset shape:", df.shape)
```

Test mode: using smaller dataset shape: (50000, 80)

This code temporarily configures pandas to display all columns so the first rows of the dataset can be fully inspected, prints them, and then restores the display setting to its default limit of 20 columns.

```
[9]: # print the first lines with all columns
pd.set_option('display.max_columns', None)
print(df.head())
pd.set_option('display.max_columns', 20) # reset to default
```

	Destination Port	Flow Duration	Total Fwd Packets	\
746827	50545	232	1	
946912	53	31226	2	
2216843	80	99951883	9	

699389	53	30894	4
1170268	53	48943	2

	Total Backward Packets	Total Length of Fwd Packets	\
746827	1	0	
946912	2	68	
2216843	7	317	
699389	2	140	
1170268	2	88	

	Total Length of Bwd Packets	Fwd Packet Length Max	\
746827	0	0	
946912	380	34	
2216843	11595	317	
699389	172	35	
1170268	166	44	

	Fwd Packet Length Min	Fwd Packet Length Mean	\
746827	0	0.000000	
946912	34	34.000000	
2216843	0	35.222222	
699389	35	35.000000	
1170268	44	44.000000	

	Fwd Packet Length Std	Bwd Packet Length Max	\
746827	0.000000	0	
946912	0.000000	190	
2216843	105.666667	5792	
699389	0.000000	86	
1170268	0.000000	83	

	Bwd Packet Length Min	Bwd Packet Length Mean	\
746827	0	0.000000	
946912	190	190.000000	
2216843	0	1656.428571	
699389	86	86.000000	
1170268	83	83.000000	

	Bwd Packet Length Std	Flow Bytes/s	Flow Packets/s	\
746827	0.000000	0.000000	8620.689655	
946912	0.000000	14347.018510	128.098380	
2216843	2118.227235	119.177345	0.160077	
699389	0.000000	10099.048360	194.212468	
1170268	0.000000	5189.710480	81.727724	

	Flow IAT Mean	Flow IAT Std	Flow IAT Max	Flow IAT Min	\
746827	2.320000e+02	0.000000e+00	232	232	
946912	1.040867e+04	1.802228e+04	31219	3	

2216843	6.663459e+06	2.580000e+07	99900000	1
699389	6.178800e+03	8.507391e+03	17161	1
1170268	1.631433e+04	2.825206e+04	48937	3

	Fwd IAT Total	Fwd IAT Mean	Fwd IAT Std	Fwd IAT Max	\
746827	0	0.0	0.000000e+00	0	
946912	3	3.0	0.000000e+00	3	
2216843	99900000	12500000.0	3.530000e+07	99900000	
699389	17211	5737.0	9.893503e+03	17161	
1170268	3	3.0	0.000000e+00	3	

	Fwd IAT Min	Bwd IAT Total	Bwd IAT Mean	Bwd IAT Std	\
746827	0	0	0.0	0.0000	
946912	3	4	4.0	0.0000	
2216843	1	56541	9423.5	18378.1771	
699389	1	49	49.0	0.0000	
1170268	3	3	3.0	0.0000	

	Bwd IAT Max	Bwd IAT Min	Fwd PSH Flags	Bwd PSH Flags	\
746827	0	0	0	0	
946912	4	4	0	0	
2216843	46050	14	0	0	
699389	49	49	0	0	
1170268	3	3	0	0	

	Fwd URG Flags	Bwd URG Flags	Fwd Header Length	\
746827	0	0	32	
946912	0	0	64	
2216843	0	0	296	
699389	0	0	128	
1170268	0	0	40	

	Bwd Header Length	Fwd Packets/s	Bwd Packets/s	\
746827	32	4310.344828	4310.344828	
946912	40	64.049190	64.049190	
2216843	232	0.090043	0.070034	
699389	64	129.474979	64.737489	
1170268	64	40.863862	40.863862	

	Min Packet Length	Max Packet Length	Packet Length Mean	\
746827	0	0	0.000000	
946912	34	190	96.400000	
2216843	0	5792	700.705882	
699389	35	86	49.571429	
1170268	44	83	59.600000	

	Packet Length Std	Packet Length Variance	FIN Flag Count	\
746827	0.000000	0.000000e+00	0	

946912	85.444719	7.300800e+03	0
2216843	1538.694445	2.367581e+06	0
699389	24.885452	6.192857e+02	0
1170268	21.361180	4.563000e+02	0
	SYN Flag Count	RST Flag Count	PSH Flag Count
746827	0	0	0
946912	0	0	0
2216843	0	0	0
699389	0	0	0
1170268	0	0	0
	URG Flag Count	CWE Flag Count	ECE Flag Count
746827	1	0	0
946912	0	0	0
2216843	0	0	0
699389	0	0	0
1170268	0	0	0
	Average Packet Size	Avg Fwd Segment Size	Avg Bwd Segment Size
746827	0.000000	0.000000	0.000000
946912	120.500000	34.000000	190.000000
2216843	744.500000	35.222222	1656.428571
699389	57.833333	35.000000	86.000000
1170268	74.500000	44.000000	83.000000
	Fwd Header Length.1	Fwd Avg Bytes/Bulk	Fwd Avg Packets/Bulk
746827	32	0	0
946912	64	0	0
2216843	296	0	0
699389	128	0	0
1170268	40	0	0
	Fwd Avg Bulk Rate	Bwd Avg Bytes/Bulk	Bwd Avg Packets/Bulk
746827	0	0	0
946912	0	0	0
2216843	0	0	0
699389	0	0	0
1170268	0	0	0
	Bwd Avg Bulk Rate	Subflow Fwd Packets	Subflow Fwd Bytes
746827	0	1	0
946912	0	2	68
2216843	0	9	317
699389	0	4	140
1170268	0	2	88
	Subflow Bwd Packets	Subflow Bwd Bytes	Init_Win_bytes_forward

746827	1	0	357		
946912	2	380	-1		
2216843	7	11595	274		
699389	2	172	-1		
1170268	2	166	-1		
	Init_Win_bytes_backward	act_data_pkt_fwd	min_seg_size_forward \		
746827	32832	0	32		
946912	-1	1	32		
2216843	235	1	32		
699389	-1	3	32		
1170268	-1	1	20		
	Active Mean	Active Std	Active Max	Active Min	Idle Mean \
746827	0.0	0.0	0	0	0.0
946912	0.0	0.0	0	0	0.0
2216843	999.0	0.0	999	999	99900000.0
699389	0.0	0.0	0	0	0.0
1170268	0.0	0.0	0	0	0.0
	Idle Std	Idle Max	Idle Min	Label	source_file
746827	0.0	0	0	BENIGN	3
946912	0.0	0	0	BENIGN	3
2216843	0.0	99900000	99900000	DoS Hulk	7
699389	0.0	0	0	BENIGN	2
1170268	0.0	0	0	BENIGN	3

This code retrieves and displays all unique values present in the 'Label' column, allowing you to see which distinct classes or categories exist in the dataset.

```
[10]: df['Label'].unique()
```

```
[10]: array(['BENIGN', 'DoS Hulk', 'DDoS', 'PortScan', 'DoS slowloris',
       'DoS GoldenEye', 'FTP-Patator', 'DoS Slowhttptest', 'Bot',
       'SSH-Patator', 'Web Attack Brute Force', 'Web Attack XSS'],
      dtype=object)
```

This code provides an overview of the dataset by displaying structural information, counting missing values per column and in total, and presenting descriptive statistics for all features to help understand the dataset's composition and potential issues.

```
[11]: # view some basic info about the dataset
print("\nDataset Info:")
print(df.info())

print("\nMissing Values per Column:")
print(df.isnull().sum())
print(f"Total number of missing values: {df.isnull().sum().sum()}")
```

```

print("\nBasic Statistics:")
display(df.describe(include='all'))

```

Dataset Info:

<class 'pandas.core.frame.DataFrame'>
Index: 50000 entries, 746827 to 2337454

Data columns (total 80 columns):

#	Column	Non-Null Count	Dtype
0	Destination Port	50000	non-null
1	Flow Duration	50000	non-null
2	Total Fwd Packets	50000	non-null
3	Total Backward Packets	50000	non-null
4	Total Length of Fwd Packets	50000	non-null
5	Total Length of Bwd Packets	50000	non-null
6	Fwd Packet Length Max	50000	non-null
7	Fwd Packet Length Min	50000	non-null
8	Fwd Packet Length Mean	50000	non-null
9	Fwd Packet Length Std	50000	non-null
10	Bwd Packet Length Max	50000	non-null
11	Bwd Packet Length Min	50000	non-null
12	Bwd Packet Length Mean	50000	non-null
13	Bwd Packet Length Std	50000	non-null
14	Flow Bytes/s	49975	non-null
15	Flow Packets/s	50000	non-null
16	Flow IAT Mean	50000	non-null
17	Flow IAT Std	50000	non-null
18	Flow IAT Max	50000	non-null
19	Flow IAT Min	50000	non-null
20	Fwd IAT Total	50000	non-null
21	Fwd IAT Mean	50000	non-null
22	Fwd IAT Std	50000	non-null
23	Fwd IAT Max	50000	non-null
24	Fwd IAT Min	50000	non-null
25	Bwd IAT Total	50000	non-null
26	Bwd IAT Mean	50000	non-null
27	Bwd IAT Std	50000	non-null
28	Bwd IAT Max	50000	non-null
29	Bwd IAT Min	50000	non-null
30	Fwd PSH Flags	50000	non-null
31	Bwd PSH Flags	50000	non-null
32	Fwd URG Flags	50000	non-null
33	Bwd URG Flags	50000	non-null
34	Fwd Header Length	50000	non-null
35	Bwd Header Length	50000	non-null
36	Fwd Packets/s	50000	non-null
37	Bwd Packets/s	50000	non-null

```

38 Min Packet Length           50000 non-null int64
39 Max Packet Length          50000 non-null int64
40 Packet Length Mean         50000 non-null float64
41 Packet Length Std          50000 non-null float64
42 Packet Length Variance     50000 non-null float64
43 FIN Flag Count             50000 non-null int64
44 SYN Flag Count             50000 non-null int64
45 RST Flag Count             50000 non-null int64
46 PSH Flag Count             50000 non-null int64
47 ACK Flag Count             50000 non-null int64
48 URG Flag Count             50000 non-null int64
49 CWE Flag Count             50000 non-null int64
50 ECE Flag Count             50000 non-null int64
51 Down/Up Ratio              50000 non-null int64
52 Average Packet Size        50000 non-null float64
53 Avg Fwd Segment Size       50000 non-null float64
54 Avg Bwd Segment Size       50000 non-null float64
55 Fwd Header Length.1        50000 non-null int64
56 Fwd Avg Bytes/Bulk        50000 non-null int64
57 Fwd Avg Packets/Bulk      50000 non-null int64
58 Fwd Avg Bulk Rate          50000 non-null int64
59 Bwd Avg Bytes/Bulk        50000 non-null int64
60 Bwd Avg Packets/Bulk      50000 non-null int64
61 Bwd Avg Bulk Rate          50000 non-null int64
62 Subflow Fwd Packets        50000 non-null int64
63 Subflow Fwd Bytes          50000 non-null int64
64 Subflow Bwd Packets        50000 non-null int64
65 Subflow Bwd Bytes          50000 non-null int64
66 Init_Win_bytes_forward     50000 non-null int64
67 Init_Win_bytes_backward     50000 non-null int64
68 act_data_pkt_fwd            50000 non-null int64
69 min_seg_size_forward       50000 non-null int64
70 Active Mean                50000 non-null float64
71 Active Std                 50000 non-null float64
72 Active Max                 50000 non-null int64
73 Active Min                 50000 non-null int64
74 Idle Mean                  50000 non-null float64
75 Idle Std                   50000 non-null float64
76 Idle Max                   50000 non-null int64
77 Idle Min                   50000 non-null int64
78 Label                       50000 non-null object
79 source_file                  50000 non-null int64

```

dtypes: float64(24), int64(55), object(1)
memory usage: 30.9+ MB
None

Missing Values per Column:

Destination Port	0
------------------	---

```

Flow Duration          0
Total Fwd Packets    0
Total Backward Packets 0
Total Length of Fwd Packets 0
..
Idle Std              0
Idle Max              0
Idle Min              0
Label                 0
source_file           0
Length: 80, dtype: int64
Total number of missing values: 25

```

Basic Statistics:

```

c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\pandas\core\nanops.py:1016: RuntimeWarning: invalid value
encountered in subtract
    sqr = _ensure_numeric((avg - values) ** 2)
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\pandas\core\nanops.py:1016: RuntimeWarning: invalid value
encountered in subtract
    sqr = _ensure_numeric((avg - values) ** 2)

```

	Destination Port	Flow Duration	Total Fwd Packets \
count	50000.000000	5.000000e+04	50000.000000
unique	Nan	Nan	Nan
top	Nan	Nan	Nan
freq	Nan	Nan	Nan
mean	8122.670040	1.467833e+07	6.404040
std	18355.671208	3.356638e+07	63.476262
min	0.000000	-1.000000e+00	1.000000
25%	53.000000	1.540000e+02	2.000000
50%	80.000000	3.128800e+04	2.000000
75%	443.000000	2.935594e+06	5.000000
max	65529.000000	1.199994e+08	9449.000000

	Total Backward Packets	Total Length of Fwd Packets \
count	50000.000000	50000.000000
unique	Nan	Nan
top	Nan	Nan
freq	Nan	Nan
mean	6.329980	518.673680
std	76.536092	2723.044996
min	0.000000	0.000000
25%	1.000000	12.000000
50%	2.000000	62.000000
75%	4.000000	193.000000
max	10063.000000	153145.000000

	Total Length of Bwd Packets	Fwd Packet Length Max	\
count	5.000000e+04	50000.000000	
unique	NaN	NaN	
top	NaN	NaN	
freq	NaN	NaN	
mean	6.986600e+03	207.416480	
std	1.650067e+05	707.720042	
min	0.000000e+00	0.000000	
25%	4.000000e+00	6.000000	
50%	1.230000e+02	37.000000	
75%	4.840000e+02	85.000000	
max	2.340090e+07	23360.000000	

	Fwd Packet Length Min	Fwd Packet Length Mean	\
count	50000.000000	50000.000000	
unique	NaN	NaN	
top	NaN	NaN	
freq	NaN	NaN	
mean	18.604720	58.313406	
std	57.708775	186.253930	
min	0.000000	0.000000	
25%	0.000000	6.000000	
50%	2.000000	34.000000	
75%	36.000000	50.000000	
max	1983.000000	4672.000000	

	Fwd Packet Length Std	...	Active Mean	Active Std	Active Max	\
count	50000.000000	...	5.000000e+04	5.000000e+04	5.000000e+04	
unique	NaN	...	NaN	NaN	NaN	
top	NaN	...	NaN	NaN	NaN	
freq	NaN	...	NaN	NaN	NaN	
mean	68.733582	...	7.749996e+04	4.327377e+04	1.535587e+05	
std	277.492447	...	5.408293e+05	4.194620e+05	1.012230e+06	
min	0.000000	...	0.000000e+00	0.000000e+00	0.000000e+00	
25%	0.000000	...	0.000000e+00	0.000000e+00	0.000000e+00	
50%	0.000000	...	0.000000e+00	0.000000e+00	0.000000e+00	
75%	26.162951	...	0.000000e+00	0.000000e+00	0.000000e+00	
max	6429.190773	...	2.920000e+07	5.040000e+07	8.740000e+07	

	Active Min	Idle Mean	Idle Std	Idle Max	Idle Min	\
count	5.000000e+04	5.000000e+04	5.000000e+04	5.000000e+04	5.000000e+04	
unique	NaN	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	NaN	
mean	5.316534e+04	8.217208e+06	4.902213e+05	8.587389e+06	7.830262e+06	
std	4.593371e+05	2.350088e+07	4.552573e+06	2.422544e+07	2.323902e+07	
min	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	

```

25%      0.000000e+00  0.000000e+00  0.000000e+00  0.000000e+00  0.000000e+00
50%      0.000000e+00  0.000000e+00  0.000000e+00  0.000000e+00  0.000000e+00
75%      0.000000e+00  0.000000e+00  0.000000e+00  0.000000e+00  0.000000e+00
max     1.500000e+07  1.200000e+08  6.680000e+07  1.200000e+08  1.200000e+08

          Label  source_file
count    50000  50000.000000
unique   12      NaN
top      BENIGN  NaN
freq    40193  NaN
mean     NaN     4.156480
std      NaN     2.348877
min      NaN     0.000000
25%     NaN     2.000000
50%     NaN     4.000000
75%     NaN     6.000000
max     NaN     7.000000

[11 rows x 80 columns]

```

This code scans all columns and drops any column that contains only one unique value, since such features provide no useful information for machine learning and can be safely removed.

```
[12]: # remove columns with the same value for all rows
for col in df.columns:
    if df[col].nunique() == 1:
        df.drop(columns=[col], inplace=True)
        print(f"Dropped column {col} with a single unique value.")
```

```

Dropped column Bwd PSH Flags with a single unique value.
Dropped column Bwd URG Flags with a single unique value.
Dropped column Fwd Avg Bytes/Bulk with a single unique value.
Dropped column Fwd Avg Packets/Bulk with a single unique value.
Dropped column Fwd Avg Bulk Rate with a single unique value.
Dropped column Bwd Avg Bytes/Bulk with a single unique value.
Dropped column Bwd Avg Packets/Bulk with a single unique value.
Dropped column Bwd Avg Bulk Rate with a single unique value.

```

This code identifies all numeric columns and all categorical (object-type) columns in the dataset, prints how many numeric features exist, and lists the names of all categorical features to prepare for later preprocessing steps.

```
[13]: numeric_cols = df.select_dtypes(include=np.number).columns
print("\nNumber of numeric Columns:", len(numeric_cols))

categorical_cols = df.select_dtypes(include=['object']).columns
print("\nCategorical Columns:", list(categorical_cols))
```

Number of numeric Columns: 71

```
Categorical Columns: [' Label']
```

This code displays all unique values found in the 'PSH Flag Count' column to help understand the range or categories present in that feature.

```
[14]: df[' PSH Flag Count'].unique()
```

```
[14]: array([0, 1])
```

Some columns contain integers, but are actually binary values 0 or 1. We will encode them as categorical variables.

```
[15]: # add columns containing only 0 and 1 to categorical columns
binary_cols = [col for col in numeric_cols if df[col].nunique() == 2]
print("\nBinary Columns (0/1):", binary_cols)
categorical_cols = binary_cols + ['source_file'] + categorical_cols.tolist()
print("\nUpdated number of Categorical Columns:", len(categorical_cols))
numeric_cols = [col for col in numeric_cols if col not in binary_cols]
print("\nUpdated number of Numeric Columns:", len(numeric_cols))
```

```
Binary Columns (0/1): ['Fwd PSH Flags', ' Fwd URG Flags', 'FIN Flag Count', ' SYN Flag Count', ' RST Flag Count', ' PSH Flag Count', ' ACK Flag Count', ' URG Flag Count', ' CWE Flag Count', ' ECE Flag Count']
```

```
Updated number of Categorical Columns: 12
```

```
Updated number of Numeric Columns: 61
```

3.1.1 Handling missing values

```
[16]: # Handle infinite or very large values before scaling
print(f"Number of infinite values before cleaning: {np.isinf(df[numeric_cols]).sum().sum()}")
df[numeric_cols] = df[numeric_cols].replace([np.inf, -np.inf], np.nan)

# Fill remaining missing values
for col in df.columns:
    if df[col].dtype == 'object':
        # categorical column, fill with most frequent value
        df[col] = df[col].fillna(df[col].mode()[0])
    else:
        # numerical column, fill with median value
        df[col] = df[col].fillna(df[col].median())
```

```
Number of infinite values before cleaning: 83
```

This code replaces infinite values in numeric features with NaN, then fills all remaining missing

values using the most frequent value for categorical columns and the median for numerical ones to ensure the dataset contains no invalid or empty entries before scaling or modeling.

3.1.2 Normalizing or scaling numerical features

The code clips extreme numerical values to prevent overflow and applies Standard Scaling to standardize features (mean=0, std=1) for model training.

```
[17]: # Clip extremely large values to avoid numerical overflow
df[numeric_cols] = df[numeric_cols].clip(lower=-1e10, upper=1e10)

scaler = StandardScaler() # to get mean = 0, std = 1
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])
```

3.1.3 Encoding categorical variables

This code uses LabelEncoder to convert each column in categorical_cols into numerical integers, ensuring string conversion first. It also prints the encoded classes for verification.

```
[18]: le = LabelEncoder()
for col in categorical_cols:
    df[col] = le.fit_transform(df[col].astype(str))
    print(f"Encoded {col} with classes: {le.classes_}")
```

```
Encoded Fwd PSH Flags with classes: ['0' '1']
Encoded Fwd URG Flags with classes: ['0' '1']
Encoded FIN Flag Count with classes: ['0' '1']
Encoded SYN Flag Count with classes: ['0' '1']
Encoded RST Flag Count with classes: ['0' '1']
Encoded PSH Flag Count with classes: ['0' '1']
Encoded ACK Flag Count with classes: ['0' '1']
Encoded URG Flag Count with classes: ['0' '1']
Encoded CWE Flag Count with classes: ['0' '1']
Encoded ECE Flag Count with classes: ['0' '1']
Encoded source_file with classes: ['-0.0666197284514735' '-0.49235930188880367'
'-0.9180988753261339'
'-1.343838448763464' '-1.7695780222007942' '0.3591198449858567'
'0.7848594184231868' '1.2105989918605171']
Encoded Label with classes: ['BENIGN' 'Bot' 'DDoS' 'DoS GoldenEye' 'DoS Hulk'
'DoS Slowhttptest'
'DoS slowloris' 'FTP-Patator' 'PortScan' 'SSH-Patator'
'Web Attack Brute Force' 'Web Attack XSS']
```


3.2 EDA - Exploratory Data Analysis

Conduct EDA, visualizing feature distributions and identifying potential relationships to guide feature engineering.

3.2.1 Visualize distribution of target variable

This block visualizes the distribution of the target variable (the ‘Label’ column) using a count plot. It retrieves the original class names using `le.inverse_transform` to correctly label the x-axis for improved readability of the class distribution.

```
[19]: target_col = 'Label'
plt.figure(figsize=(8, 4))
sns.countplot(x=df[target_col])
plt.title("Class Distribution")
class_names = le.inverse_transform(sorted(df[target_col].unique()))
plt.xticks(ticks=range(len(class_names)), labels=class_names, rotation=45)
plt.show()
```




Benign is the most frequent class. Frequent attacks are DDoS, DoS Hulk and PortScan.

3.2.2 Visualize feature distributions

This block generates histograms with 30 bins for all numerical features in `numeric_cols` across a large figure (20x20) to visualize their distributions, using `tight_layout()` to prevent overlap.

```
[20]: df[numeric_cols].hist(bins=30, figsize=(20, 20))
plt.suptitle("Feature Distributions", fontsize=18)
plt.tight_layout()
```

```
plt.show()
```



3.2.3 Correlation analysis

This block calculates the correlation matrix and performs feature selection by dropping columns that have a near-zero absolute correlation (< 0.01) with the target ‘Label’. It then prints the resulting DataFrame’s head and information, and visualizes the full correlation matrix using a heatmap.

```
[21]: corr = df.corr()
df = df.drop(corr["Label"] [abs(corr["Label"])<0.01].index, axis=1)
print(df.head(), df.info())
plt.figure(figsize=(14, 10))
```

```

sns.heatmap(corr, cmap="coolwarm", center=0)
plt.title("Feature Correlation Heatmap")
plt.show()

```

```

<class 'pandas.core.frame.DataFrame'>
Index: 50000 entries, 746827 to 2337454
Data columns (total 61 columns):
 #   Column           Non-Null Count Dtype  
 ---  -- 
 0   Destination Port    50000 non-null float64
 1   Flow Duration      50000 non-null float64
 2   Total Fwd Packets  50000 non-null float64
 3   Total Backward Packets 50000 non-null float64
 4   Total Length of Fwd Packets 50000 non-null float64
 5   Fwd Packet Length Max 50000 non-null float64
 6   Fwd Packet Length Min 50000 non-null float64
 7   Fwd Packet Length Mean 50000 non-null float64
 8   Fwd Packet Length Std 50000 non-null float64
 9   Bwd Packet Length Max 50000 non-null float64
 10  Bwd Packet Length Min 50000 non-null float64
 11  Bwd Packet Length Mean 50000 non-null float64
 12  Bwd Packet Length Std 50000 non-null float64
 13  Flow Bytes/s       50000 non-null float64
 14  Flow Packets/s     50000 non-null float64
 15  Flow IAT Mean      50000 non-null float64
 16  Flow IAT Std       50000 non-null float64
 17  Flow IAT Max       50000 non-null float64
 18  Fwd IAT Total      50000 non-null float64
 19  Fwd IAT Mean       50000 non-null float64
 20  Fwd IAT Std       50000 non-null float64
 21  Fwd IAT Max       50000 non-null float64
 22  Fwd IAT Min       50000 non-null float64
 23  Bwd IAT Total      50000 non-null float64
 24  Bwd IAT Std       50000 non-null float64
 25  Bwd IAT Max       50000 non-null float64
 26  Bwd IAT Min       50000 non-null float64
 27  Fwd PSH Flags      50000 non-null int64 
 28  Fwd Header Length  50000 non-null float64
 29  Bwd Header Length  50000 non-null float64
 30  Fwd Packets/s     50000 non-null float64
 31  Bwd Packets/s     50000 non-null float64
 32  Min Packet Length 50000 non-null float64
 33  Max Packet Length 50000 non-null float64
 34  Packet Length Mean 50000 non-null float64
 35  Packet Length Std 50000 non-null float64
 36  Packet Length Variance 50000 non-null float64
 37  FIN Flag Count     50000 non-null int64 
 38  SYN Flag Count     50000 non-null int64 

```

```

39 PSH Flag Count           50000 non-null int64
40 URG Flag Count           50000 non-null int64
41 Down/Up Ratio             50000 non-null float64
42 Average Packet Size       50000 non-null float64
43 Avg Fwd Segment Size      50000 non-null float64
44 Avg Bwd Segment Size      50000 non-null float64
45 Fwd Header Length.1       50000 non-null float64
46 Subflow Fwd Packets       50000 non-null float64
47 Subflow Fwd Bytes          50000 non-null float64
48 Subflow Bwd Packets       50000 non-null float64
49 Init_Win_bytes_forward     50000 non-null float64
50 Init_Win_bytes_backward     50000 non-null float64
51 act_data_pkt_fwd           50000 non-null float64
52 min_seg_size_forward       50000 non-null float64
53 Active Std                  50000 non-null float64
54 Active Max                  50000 non-null float64
55 Idle Mean                   50000 non-null float64
56 Idle Std                     50000 non-null float64
57 Idle Max                     50000 non-null float64
58 Idle Min                     50000 non-null float64
59 Label                         50000 non-null int64
60 source_file                   50000 non-null int64

dtypes: float64(54), int64(7)
memory usage: 23.7 MB

      Destination Port    Flow Duration   Total Fwd Packets \
746827              2.311152      -0.437290      -0.085136
946912              -0.439633      -0.436367      -0.069382
2216843             -0.438162      2.540471       0.040897
699389              -0.439633      -0.436376      -0.037873
1170268             -0.439633      -0.435839      -0.069382

      Total Backward Packets  Total Length of Fwd Packets \
746827                  -0.069641      -0.190478
946912                  -0.056575      -0.165505
2216843                  0.008754      -0.074063
699389                  -0.056575      -0.139064
1170268                  -0.056575      -0.158160


      Fwd Packet Length Max   Fwd Packet Length Min \
746827                  -0.293080      -0.322393
946912                  -0.245038      0.266778
2216843                  0.154842      -0.322393
699389                  -0.243625      0.284107
1170268                  -0.230908      0.440064

      Fwd Packet Length Mean  Fwd Packet Length Std \
746827                  -0.313089      -0.247698
946912                  -0.130540      -0.247698

```

2216843		-0.123978		0.133097			
699389		-0.125171		-0.247698			
1170268		-0.076850		-0.247698			
	Bwd Packet Length Max	...	act_data_pkt_fwd	min_seg_size_forward	\		
746827		-0.447062	...	-0.069280		0.896237	
946912		-0.348742	...	-0.048303		0.896237	
2216843		2.550154	...	-0.048303		0.896237	
699389		-0.402559	...	-0.006350		0.896237	
1170268		-0.404112	...	-0.048303		-0.923312	
	Active Std	Active Max	Idle Mean	Idle Std	Idle Max	Idle Min	\
746827	-0.103166	-0.151705	-0.349659	-0.107681	-0.354482	-0.336948	
946912	-0.103166	-0.151705	-0.349659	-0.107681	-0.354482	-0.336948	
2216843	-0.103166	-0.150718	3.901289	-0.107681	3.769324	3.961900	
699389	-0.103166	-0.151705	-0.349659	-0.107681	-0.354482	-0.336948	
1170268	-0.103166	-0.151705	-0.349659	-0.107681	-0.354482	-0.336948	
	Label	source_file					
746827	0	1					
946912	0	1					
2216843	4	7					
699389	0	2					
1170268	0	1					

[5 rows x 61 columns] None



The entirely white rows and columns in the heatmap indicate features that contain only a single value, resulting in zero variance and no correlation with other features.

The heatmap analysis indicates strong correlations between some features (e.g., Total Fwd/Backward Packets), suggesting redundancy that dimensionality reduction (like PCA) could address. It also confirms that other features are strongly correlated with the target ('Label') (e.g., PSH Flag Count, Min Packet Length), validating their importance for the classification task.

3.2.4 Create train and test datasets, and balance the dataset with SMOTE

This block splits the data into training (80%) and testing (20%) sets using a fixed random state. It then addresses class imbalance by applying the SMOTE (Synthetic Minority Over-sampling Technique) algorithm to the training data (X_{train} , y_{train}), generating synthetic samples for minority classes to achieve a balanced distribution. The code prints the sizes of the sets and the class distributions before and after SMOTE

```
[22]: X = df.drop(columns=[target_col])
y = df[target_col]
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=42)
print("Training set size:", X_train.shape)
print("Test set size:", X_test.shape)

print("Class répartition before SMOTE :")
print(y_train.value_counts())

# Appliquer SMOTE pour équilibrer les classes sur les données d'entraînement
smote = SMOTE(sampling_strategy='auto', random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

print("\nClass répartition after SMOTE :")
print(pd.Series(y_train_smote).value_counts())
print(X_train.info(), df.info())

```

Training set size: (40000, 60)

Test set size: (10000, 60)

Class répartition before SMOTE :

Label	Count
0	32125
4	3266
8	2287
2	1785
3	142
7	107
6	86
9	77
5	75
1	23
10	16
11	11

Name: count, dtype: int64

Class répartition after SMOTE :

Label	Count
2	32125
0	32125
4	32125
8	32125
5	32125
7	32125
3	32125
9	32125
6	32125
1	32125
10	32125

```

11    32125
Name: count, dtype: int64
<class 'pandas.core.frame.DataFrame'>
Index: 40000 entries, 182468 to 895607
Data columns (total 60 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Destination Port    40000 non-null   float64
 1   Flow Duration       40000 non-null   float64
 2   Total Fwd Packets  40000 non-null   float64
 3   Total Backward Packets 40000 non-null   float64
 4   Total Length of Fwd Packets 40000 non-null   float64
 5   Fwd Packet Length Max 40000 non-null   float64
 6   Fwd Packet Length Min 40000 non-null   float64
 7   Fwd Packet Length Mean 40000 non-null   float64
 8   Fwd Packet Length Std 40000 non-null   float64
 9   Bwd Packet Length Max 40000 non-null   float64
 10  Bwd Packet Length Min 40000 non-null   float64
 11  Bwd Packet Length Mean 40000 non-null   float64
 12  Bwd Packet Length Std 40000 non-null   float64
 13  Flow Bytes/s        40000 non-null   float64
 14  Flow Packets/s      40000 non-null   float64
 15  Flow IAT Mean       40000 non-null   float64
 16  Flow IAT Std        40000 non-null   float64
 17  Flow IAT Max        40000 non-null   float64
 18  Fwd IAT Total       40000 non-null   float64
 19  Fwd IAT Mean        40000 non-null   float64
 20  Fwd IAT Std         40000 non-null   float64
 21  Fwd IAT Max         40000 non-null   float64
 22  Fwd IAT Min         40000 non-null   float64
 23  Bwd IAT Total       40000 non-null   float64
 24  Bwd IAT Std         40000 non-null   float64
 25  Bwd IAT Max         40000 non-null   float64
 26  Bwd IAT Min         40000 non-null   float64
 27  Fwd PSH Flags       40000 non-null   int64 
 28  Fwd Header Length   40000 non-null   float64
 29  Bwd Header Length   40000 non-null   float64
 30  Fwd Packets/s       40000 non-null   float64
 31  Bwd Packets/s       40000 non-null   float64
 32  Min Packet Length   40000 non-null   float64
 33  Max Packet Length   40000 non-null   float64
 34  Packet Length Mean  40000 non-null   float64
 35  Packet Length Std   40000 non-null   float64
 36  Packet Length Variance 40000 non-null   float64
 37  FIN Flag Count      40000 non-null   int64 
 38  SYN Flag Count      40000 non-null   int64 
 39  PSH Flag Count      40000 non-null   int64 
 40  URG Flag Count      40000 non-null   int64

```

```

41 Down/Up Ratio           40000 non-null float64
42 Average Packet Size    40000 non-null float64
43 Avg Fwd Segment Size   40000 non-null float64
44 Avg Bwd Segment Size   40000 non-null float64
45 Fwd Header Length.1    40000 non-null float64
46 Subflow Fwd Packets    40000 non-null float64
47 Subflow Fwd Bytes      40000 non-null float64
48 Subflow Bwd Packets    40000 non-null float64
49 Init_Win_bytes_forward 40000 non-null float64
50 Init_Win_bytes_backward 40000 non-null float64
51 act_data_pkt_fwd       40000 non-null float64
52 min_seg_size_forward   40000 non-null float64
53 Active Std              40000 non-null float64
54 Active Max              40000 non-null float64
55 Idle Mean               40000 non-null float64
56 Idle Std                40000 non-null float64
57 Idle Max                40000 non-null float64
58 Idle Min                40000 non-null float64
59 source_file             40000 non-null int64

dtypes: float64(54), int64(6)
memory usage: 18.6 MB
<class 'pandas.core.frame.DataFrame'>
Index: 50000 entries, 746827 to 2337454
Data columns (total 61 columns):
 #  Column                      Non-Null Count Dtype
 --- 
 0  Destination Port            50000 non-null float64
 1  Flow Duration              50000 non-null float64
 2  Total Fwd Packets          50000 non-null float64
 3  Total Backward Packets     50000 non-null float64
 4  Total Length of Fwd Packets 50000 non-null float64
 5  Fwd Packet Length Max      50000 non-null float64
 6  Fwd Packet Length Min      50000 non-null float64
 7  Fwd Packet Length Mean     50000 non-null float64
 8  Fwd Packet Length Std      50000 non-null float64
 9  Bwd Packet Length Max      50000 non-null float64
 10 Bwd Packet Length Min      50000 non-null float64
 11 Bwd Packet Length Mean     50000 non-null float64
 12 Bwd Packet Length Std      50000 non-null float64
 13 Flow Bytes/s              50000 non-null float64
 14 Flow Packets/s             50000 non-null float64
 15 Flow IAT Mean              50000 non-null float64
 16 Flow IAT Std               50000 non-null float64
 17 Flow IAT Max               50000 non-null float64
 18 Fwd IAT Total              50000 non-null float64
 19 Fwd IAT Mean               50000 non-null float64
 20 Fwd IAT Std                50000 non-null float64
 21 Fwd IAT Max                50000 non-null float64

```

```

22 Fwd IAT Min           50000 non-null float64
23 Bwd IAT Total         50000 non-null float64
24 Bwd IAT Std           50000 non-null float64
25 Bwd IAT Max           50000 non-null float64
26 Bwd IAT Min           50000 non-null float64
27 Fwd PSH Flags          50000 non-null int64
28 Fwd Header Length       50000 non-null float64
29 Bwd Header Length       50000 non-null float64
30 Fwd Packets/s          50000 non-null float64
31 Bwd Packets/s          50000 non-null float64
32 Min Packet Length       50000 non-null float64
33 Max Packet Length       50000 non-null float64
34 Packet Length Mean       50000 non-null float64
35 Packet Length Std        50000 non-null float64
36 Packet Length Variance    50000 non-null float64
37 FIN Flag Count          50000 non-null int64
38 SYN Flag Count          50000 non-null int64
39 PSH Flag Count          50000 non-null int64
40 URG Flag Count          50000 non-null int64
41 Down/Up Ratio           50000 non-null float64
42 Average Packet Size       50000 non-null float64
43 Avg Fwd Segment Size      50000 non-null float64
44 Avg Bwd Segment Size      50000 non-null float64
45 Fwd Header Length.1       50000 non-null float64
46 Subflow Fwd Packets       50000 non-null float64
47 Subflow Fwd Bytes         50000 non-null float64
48 Subflow Bwd Packets       50000 non-null float64
49 Init_Win_bytes_forward     50000 non-null float64
50 Init_Win_bytes_backward     50000 non-null float64
51 act_data_pkt_fwd          50000 non-null float64
52 min_seg_size_forward       50000 non-null float64
53 Active Std                 50000 non-null float64
54 Active Max                 50000 non-null float64
55 Idle Mean                  50000 non-null float64
56 Idle Std                   50000 non-null float64
57 Idle Max                   50000 non-null float64
58 Idle Min                   50000 non-null float64
59 Label                      50000 non-null int64
60 source_file                  50000 non-null int64

dtypes: float64(54), int64(7)
memory usage: 23.7 MB
None None


```

3.2.5 Pairwise feature relationships

This block splits the data into training (80%) and testing (20%) sets using a fixed random state. It then addresses class imbalance by applying the SMOTE (Synthetic Minority Over-sampling Technique) algorithm to the training data (X_{train} , y_{train}), generating synthetic samples for

minority classes to achieve a balanced distribution. The code prints the sizes of the sets and the class distributions before and after SMOTE application.

```
[23]: sample_features = list(numeric_cols[:4]) # visualize only a few to keep plots readable and time reasonable
plot_features = sample_features + [target_col]
sns.pairplot(df[plot_features], hue=target_col, diag_kind="kde")
plt.suptitle("Pairwise Feature Relationships", y=1.02)
plt.show()
```



4 2. Model Training, Comparison, and Ensemble

4.0.1 2.1 Logistic regression

On the original train set, before SMOTE:

This block initializes and trains a Logistic Regression model on the training data. It then evaluates the model's performance on the test set, printing the raw confusion matrix and a full classification report. Finally, it visualizes the confusion matrix using a heatmap with a logarithmic scale (via np.log1p) for better visualization when classes are imbalanced.

```
[24]: logistic_regression = LogisticRegression()
logistic_regression.fit(X_train, y_train)
models.append({"model": logistic_regression, "name": "log_reg", "library": "sklearn"})

### Evaluate the model
y_pred = logistic_regression.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion matrix: ", conf_matrix)
print(classification_report(y_test, y_pred))

plt.figure(figsize=(8, 6))
# Use log scale for better visualization when one class is very frequent
conf_matrix_log = np.log1p(conf_matrix) # log(1 + x) to handle zeros
sns.heatmap(conf_matrix_log, annot=True, fmt='1f', cmap='Blues')
plt.title("Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()
```

```
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:
lbfgs failed to converge after 100 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=100). You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

`n_iter_i = _check_optimize_result(`

```
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\metrics\_classification.py:1731:
```

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\metrics\_classification.py:1731:
```

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\metrics\_classification.py:1731:
```

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
```

Confusion matrix: [[8014 0 7 0 20 2 2 6 14 3 0
0]
[[7 0 0 0 0 0 0 2 0 0 0]
[20 0 422 0 0 0 0 0 0 0 0]
[2 0 0 22 0 0 0 0 0 0 0]
[5 0 3 1 804 0 0 0 0 0 0]
[4 0 0 0 0 14 1 0 0 0 0]
[1 0 0 0 0 0 10 1 7 0 0]
[21 0 0 0 0 0 0 16 0 0 0]
[2 0 0 0 0 0 0 0 541 0 0]
[17 0 0 0 0 0 0 0 0 0 0]
[5 0 0 0 0 0 0 0 0 0 0]
[4 0 0 0 0 0 0 0 0 0 0]]

	precision	recall	f1-score	support
0	0.99	0.99	0.99	8068
1	0.00	0.00	0.00	9
2	0.98	0.95	0.97	442
3	0.96	0.92	0.94	24
4	0.98	0.99	0.98	813
5	0.88	0.74	0.80	19
6	0.77	0.53	0.62	19
7	0.70	0.43	0.53	37
8	0.96	1.00	0.98	543
9	0.00	0.00	0.00	17
10	0.00	0.00	0.00	5
11	0.00	0.00	0.00	4
accuracy			0.98	10000
macro avg	0.60	0.55	0.57	10000
weighted avg	0.98	0.98	0.98	10000



A lot of attacks are not detected by logistic regression, especially minority classes. The model tends to predict the majority class “Benign” too often. Still, we get decent f1-scores for some classes like DDoS and DoS Hulk.

On the SMOTE balanced train set:

This block trains a second Logistic Regression model, increasing max_iter to 1000 for convergence, but crucially, using the balanced training data (`X_train_smote`, `y_train_smote`) generated by SMOTE. The model is then evaluated on the original, imbalanced test set, with its performance metrics (confusion matrix and classification report) and log-scale heatmap visualized.

```
[25]: logistic_regression_smote = LogisticRegression(max_iter=1000)
logistic_regression_smote.fit(X_train_smote, y_train_smote)
models.append({"model": logistic_regression_smote, "name": "log_reg_smote", "library": "sklearn"})
```

```

### Evaluate the model
y_pred = logistic_regression_smote.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion matrix: ", conf_matrix)
print(classification_report(y_test, y_pred))

plt.figure(figsize=(8, 6))
# Use log scale for better visualization when one class is very frequent
conf_matrix_log = np.log1p(conf_matrix) # log(1 + x) to handle zeros
sns.heatmap(conf_matrix_log, annot=True, fmt='.1f', cmap='Blues')
plt.title("Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()

```

```

c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:
lbfgs failed to converge after 1000 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT

```

Increase the number of iterations to improve the convergence (max_iter=1000). You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>


Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Confusion matrix:															
		precision	recall	f1-score	support										
0	1.00	0.89	0.94	8068											
1	0	441	0	0	0	0	0	0	0	0	0	0	0		
2	7	0	0	0	0	0	0	0	0	0	0	0	0		
3	0	24	0	0	0	0	0	0	0	0	0	0	0		
4	0	3	1	809	0	0	0	0	0	0	0	0	0		
5	1	0	0	0	18	0	0	0	0	0	0	0	0		
6	0	0	0	0	1	18	0	0	0	0	0	0	0		
7	0	0	0	0	0	0	37	0	0	0	0	0	0		
8	0	0	0	0	0	0	1	541	0	0	0	1	1		
9	0	0	0	0	0	0	0	0	16	0	0	1	1		
10	0	0	0	0	0	0	0	0	2	3	0	0	0		
11	0	0	0	0	0	0	0	0	0	3	1	1	1		

1	0.01	0.78	0.02	9
2	0.83	1.00	0.91	442
3	0.67	1.00	0.80	24
4	0.97	1.00	0.98	813
5	0.60	0.95	0.73	19
6	0.38	0.95	0.55	19
7	0.60	1.00	0.75	37
8	0.97	1.00	0.99	543
9	0.25	0.94	0.40	17
10	0.17	0.60	0.26	5
11	0.02	0.25	0.04	4
accuracy			0.91	10000
macro avg	0.54	0.86	0.61	10000
weighted avg	0.98	0.91	0.94	10000



4.1 2.2 KNN implementation :

On the original train set, before SMOTE:

This block trains a K-Nearest Neighbors (KNN) classifier with $n = 5$ neighbors and ‘distance’ weighting on the original training data. It evaluates the model on the test set, displaying the confusion matrix, the classification report, and a log-scale heatmap visualization of the results.

```
[26]: knn = KNeighborsClassifier(n_neighbors=5, weights='distance')
knn.fit(X_train, y_train)
models.append({"model" : knn, "name": "knn", "library": "sklearn"})

y_pred = knn.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion matrix: ", conf_matrix)
print(classification_report(y_test, y_pred))

plt.figure(figsize=(8, 6))
# Use log scale for better visualization when one class is very frequent
conf_matrix_log = np.log1p(conf_matrix) # log(1 + x) to handle zeros
sns.heatmap(conf_matrix_log, annot=True, fmt='.1f', cmap='Blues')
plt.title("Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()
```


Confusion matrix: [[8042 2 1 3 12 0 1 2 2 0 1
2]

```
[ 4 5 0 0 0 0 0 0 0 0 0 ]  
[ 0 0 442 0 0 0 0 0 0 0 0 ]  
[ 0 0 0 24 0 0 0 0 0 0 0 ]  
[ 3 0 1 0 809 0 0 0 0 0 0 ]  
[ 1 0 0 0 0 18 0 0 0 0 0 ]  
[ 0 0 0 0 0 0 19 0 0 0 0 ]  
[ 1 0 0 0 0 0 0 36 0 0 0 ]  
[ 1 0 0 0 0 0 0 0 542 0 0 ]  
[ 1 0 0 0 0 0 0 0 0 16 0 ]  
[ 2 0 0 0 0 0 0 0 0 0 3 ]  
[ 0 0 0 0 0 0 0 0 0 0 4 ]]
```

precision recall f1-score support

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8068
1	0.71	0.56	0.62	9
2	1.00	1.00	1.00	442

3	0.89	1.00	0.94	24
4	0.99	1.00	0.99	813
5	1.00	0.95	0.97	19
6	0.95	1.00	0.97	19
7	0.95	0.97	0.96	37
8	1.00	1.00	1.00	543
9	1.00	0.94	0.97	17
10	0.38	0.60	0.46	5
11	0.00	0.00	0.00	4
accuracy			1.00	10000
macro avg	0.82	0.83	0.82	10000
weighted avg	1.00	1.00	1.00	10000



On the SMOTE balanced train set:

This block trains a second KNN model with the same hyperparameters but uses the balanced,

SMOTE-resampled training data (`X_train_smote`, `y_train_smote`). It evaluates performance on the original test set, displaying the classification report and visualizing the log-scale confusion matrix heatmap for comparison.


```
[27]: knn_smote = KNeighborsClassifier(n_neighbors=5, weights='distance')
knn_smote.fit(X_train_smote, y_train_smote)
models.append({"model": knn_smote, "name": "knn_smote", "library": "sklearn"})
y_pred = knn_smote.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion matrix: ", conf_matrix)
print(classification_report(y_test, y_pred))

plt.figure(figsize=(8, 6))
# Use log scale for better visualization when one class is very frequent
conf_matrix_log = np.log1p(conf_matrix) # log(1 + x) to handle zeros
sns.heatmap(conf_matrix_log, annot=True, fmt='.1f', cmap='Blues')
plt.title("Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()
```

```
Confusion matrix: [[8016    12     1     4    14     1     6     3     4     4     3
0]
 [ 2      7     0     0     0     0     0     0     0     0     0]
 [ 0      0  442     0     0     0     0     0     0     0     0]
 [ 0      0     0   24     0     0     0     0     0     0     0]
 [ 0      0     0     1  812     0     0     0     0     0     0]
 [ 0      0     0     0     0   19     0     0     0     0     0]
 [ 0      0     0     0     0     0   19     0     0     0     0]
 [ 0      0     0     0     0     0     0   37     0     0     0]
 [ 0      0     0     0     0     0     0     0  543     0     0]
 [ 1      0     0     0     0     0     0     0     0     0   16
0]
 [ 2      0     0     0     0     0     0     0     0     0     3
0]
 [ 0      0     0     0     0     0     0     0     0     0     3   11]]
```

	precision	recall	f1-score	support
0	1.00	0.99	1.00	8068
1	0.37	0.78	0.50	9
2	1.00	1.00	1.00	442
3	0.83	1.00	0.91	24
4	0.98	1.00	0.99	813
5	0.95	1.00	0.97	19
6	0.76	1.00	0.86	19

7	0.93	1.00	0.96	37
8	0.99	1.00	1.00	543
9	0.80	0.94	0.86	17
10	0.33	0.60	0.43	5
11	1.00	0.25	0.40	4
accuracy			0.99	10000
macro avg	0.83	0.88	0.82	10000
weighted avg	1.00	0.99	0.99	10000



4.2 2.4 Random Forest Classification

On the original train set, before SMOTE:

This block trains a Random Forest classifier with 100 estimators, using the class_weight='balanced' setting to automatically adjust weights inversely proportional to class frequencies to handle imbal-

ance. It then evaluates the model's performance on the test set, printing the results and visualizing the log-scale confusion matrix.

```
[28]: random_forest = RandomForestClassifier(n_estimators=100, random_state=42, class_weight='balanced')
random_forest.fit(X_train, y_train)
models.append({"model": random_forest, "name": "rand_forest", "library": "sklearn"})
y_pred_rf = random_forest.predict(X_test)

conf_matrix = confusion_matrix(y_test, y_pred_rf)
print("Random Forest Results:")
print("Confusion matrix: ", conf_matrix)
print(classification_report(y_test, y_pred_rf))

plt.figure(figsize=(8, 6))
conf_matrix_log = np.log1p(conf_matrix)
sns.heatmap(conf_matrix_log, annot=True, fmt='.1f', cmap='Blues')
plt.title("Random Forest Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()
```

Random Forest Results:

Confusion matrix: [[8065 0 0 1 2 0 0 0 0 0 0 0 0]

```
[ 4 5 0 0 0 0 0 0 0 0 0 0 0]
[ 0 0 442 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 24 0 0 0 0 0 0 0 0 0]
[ 2 0 0 0 811 0 0 0 0 0 0 0 0]
[ 1 0 0 0 0 17 1 0 0 0 0 0 0]
[ 0 0 0 0 0 0 19 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 37 0 0 0 0 0]
[ 1 0 0 0 0 0 0 0 0 542 0 0 0]
[ 1 0 0 0 0 0 0 0 0 0 16 0 0]
[ 3 0 0 0 0 0 0 0 0 0 0 2 0]
[ 0 0 0 0 0 0 0 0 0 0 0 4 0]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	8068
1	1.00	0.56	0.71	9
2	1.00	1.00	1.00	442
3	0.96	1.00	0.98	24
4	1.00	1.00	1.00	813


5	1.00	0.89	0.94	19
6	0.95	1.00	0.97	19
7	1.00	1.00	1.00	37
8	1.00	1.00	1.00	543
9	1.00	0.94	0.97	17
10	0.33	0.40	0.36	5
11	0.00	0.00	0.00	4
accuracy			1.00	10000
macro avg	0.85	0.82	0.83	10000
weighted avg	1.00	1.00	1.00	10000

```
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\metrics\_classification.py:1731:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\metrics\_classification.py:1731:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\metrics\_classification.py:1731:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.

    _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
```



On the SMOTE balanced train set:

This block trains a second Random Forest model, retaining the `class_weight='balanced'` parameter but crucially fitting it using the SMOTE-resampled training data. The model is then evaluated on the original test set, with its performance (classification report and log-scale confusion matrix) printed and visualized.

```
[29]: random_forest_smote = RandomForestClassifier(n_estimators=100, random_state=42, ↴
    ↴class_weight='balanced')
random_forest_smote.fit(X_train_smote, y_train_smote)
models.append({"model": random_forest_smote, "name": "rand_forest_smote", ↴
    ↴"library": "sklearn"})
y_pred_rf = random_forest_smote.predict(X_test)

conf_matrix = confusion_matrix(y_test, y_pred_rf)
print("Random Forest Results:")
print("Confusion matrix: ", conf_matrix)
```

```


print(classification_report(y_test, y_pred_rf))

plt.figure(figsize=(8, 6))
conf_matrix_log = np.log1p(conf_matrix)
sns.heatmap(conf_matrix_log, annot=True, fmt='1f', cmap='Blues')
plt.title("Random Forest Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()

```

Random Forest Results:

	0	1	2	3	4	5	6	7	8	9	10	11	accuracy	macro avg	weighted avg
Confusion matrix:	[[8062	1	0	1	2	1	0	0	0	0	0	1	0	0	0
	0]														
	[3	6	0	0	0	0	0	0	0	0	0	0	0	0	0]
	[0	0	442	0	0	0	0	0	0	0	0	0	0	0	0]
	[0	0	0	24	0	0	0	0	0	0	0	0	0	0	0]
	[2	0	0	0	811	0	0	0	0	0	0	0	0	0	0]
	[0	0	0	0	0	18	1	0	0	0	0	0	0	0	0]
	[0	0	0	0	0	0	19	0	0	0	0	0	0	0	0]
	[0	0	0	0	0	0	0	37	0	0	0	0	0	0	0]
	[0	0	0	0	0	0	0	0	543	0	0	0	0	0	0]
	[1	0	0	0	0	0	0	0	0	16	0	0	0	0	0]
	[2	0	0	0	0	0	0	0	0	0	3	0	0	0	0]
	[0	0	0	0	0	0	0	0	0	0	2	2	2	2	2]
		precision	recall	f1-score	support										
		0	1.00	1.00	1.00	8068									
		1	0.86	0.67	0.75	9									
		2	1.00	1.00	1.00	442									
		3	0.96	1.00	0.98	24									
		4	1.00	1.00	1.00	813									
		5	0.95	0.95	0.95	19									
		6	0.95	1.00	0.97	19									
		7	1.00	1.00	1.00	37									
		8	1.00	1.00	1.00	543									
		9	1.00	0.94	0.97	17									
		10	0.50	0.60	0.55	5									
		11	1.00	0.50	0.67	4									
		accuracy			1.00	10000									
		macro avg			0.90	10000									
		weighted avg			1.00	10000									



4.3 2.xg Boost Classification

On the original train set, before SMOTE:

This block trains an XGBoost classifier, utilizing the faster hist tree method and implementing early stopping (stopping after 2 rounds if performance on the test set does not improve). The model is trained on the original training data, evaluated on the test set, and its performance metrics (classification report and log-scale confusion matrix) are printed and visualized.

```
[30]: gradient_boosting = xgb.XGBClassifier(tree_method="hist",
                                         early_stopping_rounds=2)
gradient_boosting.fit(X_train, y_train, eval_set=[(X_test, y_test)])
models.append({"model": gradient_boosting, "name": "XGB", "library": "xgb"})
y_pred_gb = gradient_boosting.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred_gb)
print("Gradient Boosting Results:")
print("Confusion matrix: ", conf_matrix)
```

```

print(classification_report(y_test, y_pred_gb))

plt.figure(figsize=(8, 6))
conf_matrix_log = np.log1p(conf_matrix)
sns.heatmap(conf_matrix_log, annot=True, fmt=' .1f ', cmap='Blues')
plt.title("Gradient Boosting Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()

```

```

[0] validation_0-mlogloss:1.00574
[1] validation_0-mlogloss:0.69199
[2] validation_0-mlogloss:0.49702
[3] validation_0-mlogloss:0.36451
[4] validation_0-mlogloss:0.27057
[5] validation_0-mlogloss:0.20249
[6] validation_0-mlogloss:0.15279
[7] validation_0-mlogloss:0.11627
[8] validation_0-mlogloss:0.08944
[9] validation_0-mlogloss:0.06951
[10] validation_0-mlogloss:0.05452
[11] validation_0-mlogloss:0.04334
[12] validation_0-mlogloss:0.03460
[13] validation_0-mlogloss:0.02785
[14] validation_0-mlogloss:0.02265
[15] validation_0-mlogloss:0.01871
[16] validation_0-mlogloss:0.01561
[17] validation_0-mlogloss:0.01334
[18] validation_0-mlogloss:0.01158
[19] validation_0-mlogloss:0.01026
[20] validation_0-mlogloss:0.00927
[21] validation_0-mlogloss:0.00844
[22] validation_0-mlogloss:0.00773
[23] validation_0-mlogloss:0.00720
[24] validation_0-mlogloss:0.00680
[25] validation_0-mlogloss:0.00631
[26] validation_0-mlogloss:0.00606
[27] validation_0-mlogloss:0.00588
[28] validation_0-mlogloss:0.00578
[29] validation_0-mlogloss:0.00560
[30] validation_0-mlogloss:0.00552
[31] validation_0-mlogloss:0.00544
[32] validation_0-mlogloss:0.00540

```

```

[33] validation_0-mlogloss:0.00529
[34] validation_0-mlogloss:0.00522
[35] validation_0-mlogloss:0.00519
[36] validation_0-mlogloss:0.00521
[37] validation_0-mlogloss:0.00516
[38] validation_0-mlogloss:0.00515
[39] validation_0-mlogloss:0.00511
[40] validation_0-mlogloss:0.00508
[41] validation_0-mlogloss:0.00503
[42] validation_0-mlogloss:0.00502
[43] validation_0-mlogloss:0.00502
[44] validation_0-mlogloss:0.00498
[45] validation_0-mlogloss:0.00495
[46] validation_0-mlogloss:0.00494
[47] validation_0-mlogloss:0.00490
[48] validation_0-mlogloss:0.00491
[49] validation_0-mlogloss:0.00489
[50] validation_0-mlogloss:0.00489
[51] validation_0-mlogloss:0.00489
[52] validation_0-mlogloss:0.00488
[53] validation_0-mlogloss:0.00491
[54] validation_0-mlogloss:0.00491

```

Gradient Boosting Results:


Confusion matrix: [[8064 0 0 1 2 1 0 0 0 0 0 0]

[1	8	0	0	0	0	0	0	0	0	0]
[0	0	442	0	0	0	0	0	0	0	0]
[0	0	0	24	0	0	0	0	0	0	0]
[0	0	0	0	813	0	0	0	0	0	0]
[0	0	0	0	0	19	0	0	0	0	0]
[0	0	0	0	0	0	19	0	0	0	0]
[0	0	0	0	0	0	0	37	0	0	0]
[0	0	0	0	0	0	0	0	543	0	0]
[1	0	0	0	0	0	0	0	0	16	0]
[2	0	0	0	0	0	0	0	0	0	3]
[0	0	0	0	0	0	0	0	0	3	1]]

precision recall f1-score support

	0	1.00	1.00	1.00	8068
1	1.00	0.89	0.94	9	
2	1.00	1.00	1.00	442	
3	0.96	1.00	0.98	24	
4	1.00	1.00	1.00	813	
5	0.95	1.00	0.97	19	
6	1.00	1.00	1.00	19	
7	1.00	1.00	1.00	37	
8	1.00	1.00	1.00	543	
9	1.00	0.94	0.97	17	

10	0.50	0.60	0.55	5
11	1.00	0.25	0.40	4
accuracy			1.00	10000
macro avg	0.95	0.89	0.90	10000
weighted avg	1.00	1.00	1.00	10000



On the SMOTE balanced train set:

This block trains a second XGBoost classifier with the same parameters (hist tree method, early_stopping_rounds=2) but utilizes the SMOTE-resampled training data. It evaluates the model on the original test set, printing the full classification report and visualizing the log-scale confusion matrix to assess the impact of balancing on the model's performance.

```
[31]: gradient_boosting_smote = xgb.XGBClassifier(tree_method="hist", ↴
    early_stopping_rounds=2)
```

```

gradient_boosting_smote.fit(X_train_smote, y_train_smote, eval_set=[(X_test, y_test)])
models.append({"model" : gradient_boosting_smote, "name": "XGB_smote", "library": "xgb"})
y_pred_gb = gradient_boosting_smote.predict(X_test)

conf_matrix = confusion_matrix(y_test, y_pred_gb)
print("Gradient Boosting Results:")
print("Confusion matrix: ", conf_matrix)
print(classification_report(y_test, y_pred_gb))

plt.figure(figsize=(8, 6))
conf_matrix_log = np.log1p(conf_matrix)
sns.heatmap(conf_matrix_log, annot=True, fmt='.1f', cmap='Blues')
plt.title("Gradient Boosting Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()

```

```

[0] validation_0-mlogloss:1.24058
[1] validation_0-mlogloss:0.82534
[2] validation_0-mlogloss:0.60220
[3] validation_0-mlogloss:0.44161
[4] validation_0-mlogloss:0.33313
[5] validation_0-mlogloss:0.25540
[6] validation_0-mlogloss:0.19896
[7] validation_0-mlogloss:0.15780
[8] validation_0-mlogloss:0.12754
[9] validation_0-mlogloss:0.10024
[10] validation_0-mlogloss:0.07889
[11] validation_0-mlogloss:0.06269
[12] validation_0-mlogloss:0.05028
[13] validation_0-mlogloss:0.04115
[14] validation_0-mlogloss:0.03384
[15] validation_0-mlogloss:0.02826
[16] validation_0-mlogloss:0.02386
[17] validation_0-mlogloss:0.02000
[18] validation_0-mlogloss:0.01730
[19] validation_0-mlogloss:0.01515
[20] validation_0-mlogloss:0.01375
[21] validation_0-mlogloss:0.01225
[22] validation_0-mlogloss:0.01113
[23] validation_0-mlogloss:0.01016

```

```

[24] validation_0-mlogloss:0.00934
[25] validation_0-mlogloss:0.00888
[26] validation_0-mlogloss:0.00845
[27] validation_0-mlogloss:0.00817
[28] validation_0-mlogloss:0.00786
[29] validation_0-mlogloss:0.00758
[30] validation_0-mlogloss:0.00741
[31] validation_0-mlogloss:0.00731
[32] validation_0-mlogloss:0.00716
[33] validation_0-mlogloss:0.00706
[34] validation_0-mlogloss:0.00692
[35] validation_0-mlogloss:0.00690
[36] validation_0-mlogloss:0.00686
[37] validation_0-mlogloss:0.00678
[38] validation_0-mlogloss:0.00674
[39] validation_0-mlogloss:0.00671
[40] validation_0-mlogloss:0.00663
[41] validation_0-mlogloss:0.00650
[42] validation_0-mlogloss:0.00641
[43] validation_0-mlogloss:0.00634
[44] validation_0-mlogloss:0.00629
[45] validation_0-mlogloss:0.00624
[46] validation_0-mlogloss:0.00620
[47] validation_0-mlogloss:0.00620
[48] validation_0-mlogloss:0.00623
[49] validation_0-mlogloss:0.00624

```

Gradient Boosting Results:

Confusion matrix: [[8060 3 0 1 3 1 0 0 0 0 0]


```

[ 2    7    0    0    0    0    0    0    0    0    0    0]
[ 0    0  442    0    0    0    0    0    0    0    0    0]
[ 0    0    0   24    0    0    0    0    0    0    0    0]
[ 0    0    0    0  813    0    0    0    0    0    0    0]
[ 0    0    0    0    0   19    0    0    0    0    0    0]
[ 0    0    0    0    0    0   19    0    0    0    0    0]
[ 0    0    0    0    0    0    0   37    0    0    0    0]
[ 0    0    0    0    0    0    0    0   543    0    0    0]
[ 0    0    0    0    0    0    0    0    0    0   17    0    0]
[ 2    0    0    0    0    0    0    0    0    0    0    3    0]
[ 0    0    0    0    0    0    0    0    0    0    0    2    2]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	8068
1	0.70	0.78	0.74	9
2	1.00	1.00	1.00	442
3	0.96	1.00	0.98	24
4	1.00	1.00	1.00	813
5	0.95	1.00	0.97	19

6	1.00	1.00	1.00	19
7	1.00	1.00	1.00	37
8	1.00	1.00	1.00	543
9	1.00	1.00	1.00	17
10	0.60	0.60	0.60	5
11	1.00	0.50	0.67	4
accuracy			1.00	10000
macro avg	0.93	0.91	0.91	10000
weighted avg	1.00	1.00	1.00	10000



5 2.6 Ensemble Methods

5.1 2.6.1 Voting Classifier

This block implements an Ensemble Learning approach by creating two types of VotingClassifier (Hard and Soft), leveraging Logistic Regression, KNN, and Random Forest as base models. Both ensemble models are trained on the SMOTE-resampled data and then evaluated on the test set, with their performance and classification reports printed. Hard Voting uses majority prediction, while Soft Voting aggregates predicted probabilities for the final result.

```
[32]: # Create base models for voting ensemble
lr_voting = LogisticRegression(max_iter=1000)
knn_voting = KNeighborsClassifier(n_neighbors=5, weights='distance')
rf_voting = RandomForestClassifier(n_estimators=100, random_state=42, class_weight='balanced')

# Hard Voting Classifier
hard_voting_clf = VotingClassifier(
    estimators=[
        ('lr', lr_voting),
        ('knn', knn_voting),
        ('rf', rf_voting),
    ],
    voting='hard'
)

print("Training Hard Voting Classifier...")
hard_voting_clf.fit(X_train_smote, y_train_smote)
models.append({"model": hard_voting_clf, "name": "hard_voting", "library": "sklearn"})
y_pred_hard = hard_voting_clf.predict(X_test)

print("Hard Voting Results:")
print(classification_report(y_test, y_pred_hard))

# Soft Voting Classifier
soft_voting_clf = VotingClassifier(
    estimators=[
        ('lr', lr_voting),
        ('knn', knn_voting),
        ('rf', rf_voting),
    ],
    voting='soft'
)

print("\nTraining Soft Voting Classifier...")
```

```

soft_voting_clf.fit(X_train_smote, y_train_smote)
models.append({"model" : soft_voting_clf, "name": "soft_voting", "library": "sklearn"})

y_pred_soft = soft_voting_clf.predict(X_test)

print("Soft Voting Results:")
print(classification_report(y_test, y_pred_soft))

```

Training Hard Voting Classifier...

```
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:
lbfgs failed to converge after 1000 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=1000).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Hard Voting Results:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8068
1	0.39	0.78	0.52	9
2	1.00	1.00	1.00	442
3	0.92	1.00	0.96	24
4	0.99	1.00	0.99	813
5	1.00	1.00	1.00	19
6	0.90	1.00	0.95	19
7	0.93	1.00	0.96	37
8	1.00	1.00	1.00	543
9	0.80	0.94	0.86	17
10	0.50	0.60	0.55	5
11	1.00	0.50	0.67	4
accuracy			1.00	10000
macro avg	0.87	0.90	0.87	10000
weighted avg	1.00	1.00	1.00	10000

Training Soft Voting Classifier...

```
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:
```

```

lbfsgs failed to converge after 1000 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT

Increase the number of iterations to improve the convergence (max_iter=1000).
You might also want to scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result()

Soft Voting Results:
      precision    recall   f1-score   support
          0         1.00     1.00     1.00      8068
          1         0.39     0.78     0.52       9
          2         1.00     1.00     1.00      442
          3         0.92     1.00     0.96      24
          4         0.99     1.00     0.99      813
          5         1.00     1.00     1.00      19
          6         0.90     1.00     0.95      19
          7         0.93     1.00     0.96      37
          8         1.00     1.00     1.00      543
          9         0.80     0.94     0.86      17
         10         0.38     0.60     0.46       5
         11         1.00     0.50     0.67       4

      accuracy                           1.00      10000
   macro avg       0.86     0.90     0.86      10000
weighted avg      1.00     1.00     1.00      10000

```

This block generates a figure with two subplots to visually compare the performance of the Hard and Soft Voting ensemble classifiers. It calculates and displays the log-scale confusion matrices for both models side-by-side, using distinct colormaps ('Blues' and 'Greens') to facilitate comparison of their misclassification patterns.

```
[33]: # Visualize voting results
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Hard voting confusion matrix
conf_matrix_hard = confusion_matrix(y_test, y_pred_hard)
conf_matrix_hard_log = np.log1p(conf_matrix_hard)
sns.heatmap(conf_matrix_hard_log, annot=True, fmt='.1f', cmap='Blues', ax=axes[0])
axes[0].set_title("Hard Voting Confusion Matrix (Log Scale)")
axes[0].set_xlabel("Predicted Label")
axes[0].set_ylabel("True Label")
```


```

axes[0].set_xticks(np.arange(len(class_names))+0.5)
axes[0].set_xticklabels(class_names, rotation=45)
axes[0].set_yticks(np.arange(len(class_names))+0.5)
axes[0].set_yticklabels(class_names, rotation=45)

# Soft voting confusion matrix
conf_matrix_soft = confusion_matrix(y_test, y_pred_soft)
conf_matrix_soft_log = np.log1p(conf_matrix_soft)
sns.heatmap(conf_matrix_soft_log, annot=True, fmt=' .1f ', cmap='Greens', □
    ↵ax=axes[1])
axes[1].set_title("Soft Voting Confusion Matrix (Log Scale)")
axes[1].set_xlabel("Predicted Label")
axes[1].set_ylabel("True Label")
axes[1].set_xticks(np.arange(len(class_names))+0.5)
axes[1].set_xticklabels(class_names, rotation=45)
axes[1].set_yticks(np.arange(len(class_names))+0.5)
axes[1].set_yticklabels(class_names, rotation=45)

plt.tight_layout()
plt.show()

```



5.2 2.6.2 Stacking Classifier

This block implements a Stacking ensemble, using Logistic Regression, KNN, and Random Forest as base estimators, with a separate Random Forest as the meta-learner to combine their predictions. The model is trained on the SMOTE-resampled data, using 5-fold cross-validation to generate meta-features. The final performance on the test set is evaluated via a classification report and visualized with a log-scale confusion matrix.

```
[34]: # Create base models for stacking
base_models = [
    ('lr', LogisticRegression()),
```

```

('knn', KNeighborsClassifier(n_neighbors=5, weights='distance')),
('rf', RandomForestClassifier(n_estimators=100, random_state=42,
                             class_weight='balanced')),
]

# Use Random Forest as meta-learner
meta_learner = RandomForestClassifier(n_estimators=50, random_state=42)

# Create stacking classifier
stacking_clf = StackingClassifier(
    estimators=base_models,
    final_estimator=meta_learner,
    cv=5 # Use 5-fold cross-validation for generating meta-features
)

print("Training Stacking Classifier...")
stacking_clf.fit(X_train_smote, y_train_smote)
models.append({"model" : stacking_clf, "name": "stacking", "library": "sklearn", "name": "stacking"})
y_pred_stack = stacking_clf.predict(X_test)

print("Stacking Results:")
print(classification_report(y_test, y_pred_stack))

# Visualize stacking results
conf_matrix_stack = confusion_matrix(y_test, y_pred_stack)
plt.figure(figsize=(8, 6))
conf_matrix_stack_log = np.log1p(conf_matrix_stack)
sns.heatmap(conf_matrix_stack_log, annot=True, fmt='.1f', cmap='Oranges')
plt.title("Stacking Classifier Confusion Matrix (Log Scale)")
plt.xlabel("Predicted Label")
plt.xticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.ylabel("True Label")
plt.yticks(ticks=np.arange(len(class_names))+0.5, labels=class_names, rotation=45)
plt.show()

```

Training Stacking Classifier...

```
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:
lbfgs failed to converge after 100 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity  
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:  
lbfgs failed to converge after 100 iteration(s) (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity  
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:  
lbfgs failed to converge after 100 iteration(s) (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity  
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:  
lbfgs failed to converge after 100 iteration(s) (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity  
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:  
lbfgs failed to converge after 100 iteration(s) (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  
c:\Users\arthu\OneDrive\Efrei\I1\S5\MachineLearning\machineLearningCybersecurity  
\Lib\site-packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning:  
lbfgs failed to converge after 100 iteration(s) (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>


Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Stacking Results:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	8068
1	0.83	0.56	0.67	9
2	1.00	1.00	1.00	442
3	0.96	1.00	0.98	24
4	1.00	1.00	1.00	813
5	1.00	0.89	0.94	19
6	0.95	1.00	0.97	19
7	1.00	1.00	1.00	37
8	1.00	1.00	1.00	543
9	1.00	0.94	0.97	17
10	0.60	0.60	0.60	5
11	1.00	0.50	0.67	4
accuracy			1.00	10000
macro avg	0.94	0.87	0.90	10000
weighted avg	1.00	1.00	1.00	10000



5.3 2.7 Model Performance Comparison

This block aggregates the performance metrics (Accuracy, F1-Macro, F1-Weighted) for all trained models (base and ensemble). It creates a comparison table, sorted by F1-Weighted score, and visualizes the results using three side-by-side bar plots to compare the models across these key metrics.

```
[40]: # Get predictions from all models for comparison
models_to_compare = {
    'Logistic Regression': logistic_regression_smote,
    'KNN': knn_smote,
    'Random Forest': random_forest_smote,
    'Gradient Boosting': gradient_boosting_smote,
    'Hard Voting': hard_voting_clf,
    'Soft Voting': soft_voting_clf,
    'Stacking': stacking_clf
```

```

}

# Calculate metrics for each model
results = []
for name, model in models_to_compare.items():
    if name in ['Hard Voting', 'Soft Voting', 'Stacking']:
        # These models are already trained and we have their predictions
        if name == 'Hard Voting':
            y_pred = y_pred_hard
        elif name == 'Soft Voting':
            y_pred = y_pred_soft
        else: # Stacking
            y_pred = y_pred_stack
    else:
        y_pred = model.predict(X_test)

    accuracy = accuracy_score(y_test, y_pred)
    f1_macro = f1_score(y_test, y_pred, average='macro')
    f1_weighted = f1_score(y_test, y_pred, average='weighted')

    results.append({
        'Model': name,
        'Accuracy': accuracy,
        'F1-Score (Macro)': f1_macro,
        'F1-Score (Weighted)': f1_weighted
    })

# Create comparison DataFrame
comparison_df = pd.DataFrame(results)
comparison_df = comparison_df.sort_values('F1-Score (Weighted)', ↴
                                           ascending=False)

print("Model Performance Comparison:")
print("=" * 80)
print(comparison_df.to_string(index=False, float_format='%.4f'))

# Visualize model comparison
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

# Accuracy comparison
axes[0].bar(range(len(comparison_df)), comparison_df['Accuracy'], ↴
            color='skyblue')
axes[0].set_title('Model Accuracy Comparison')
axes[0].set_ylabel('Accuracy')
axes[0].set_xticks(range(len(comparison_df)))
axes[0].set_xticklabels(comparison_df['Model'], rotation=45, ha='right')
axes[0].set_ylim(0, 1)

```

```

# F1-Score (Macro) comparison
axes[1].bar(range(len(comparison_df)), comparison_df['F1-Score (Macro)'], color='lightgreen')
axes[1].set_title('F1-Score (Macro) Comparison')
axes[1].set_ylabel('F1-Score (Macro)')
axes[1].set_xticks(range(len(comparison_df)))
axes[1].set_xticklabels(comparison_df['Model'], rotation=45, ha='right')
axes[1].set_ylim(0, 1)


# F1-Score (Weighted) comparison
axes[2].bar(range(len(comparison_df)), comparison_df['F1-Score (Weighted)'], color='orange')
axes[2].set_title('F1-Score (Weighted) Comparison')
axes[2].set_ylabel('F1-Score (Weighted)')
axes[2].set_xticks(range(len(comparison_df)))
axes[2].set_xticklabels(comparison_df['Model'], rotation=45, ha='right')
axes[2].set_ylim(0, 1)

plt.tight_layout()
plt.show()

```

Model Performance Comparison:

Model	Accuracy	F1-Score (Macro)	F1-Score (Weighted)
Gradient Boosting	0.9986	0.9129	0.9986
Random Forest	0.9983	0.9025	0.9983
Stacking	0.9983	0.8998	0.9982
Hard Voting	0.9956	0.8712	0.9958
Soft Voting	0.9955	0.8643	0.9958
KNN	0.9939	0.8234	0.9943
Logistic Regression	0.9073	0.6133	0.9396



5.4 2.8 Feature Importance Analysis

This block calculates and compares feature importance scores from the Random Forest and Gradient Boosting models, ranking features for each. It then visualizes the top 15 most important features for both models using horizontal bar charts and prints the top 10 features in a table format.

```
[36]: # Extract feature importance from tree-based models
feature_names = X.columns

# Random Forest feature importance
rf_importance = random_forest.feature_importances_
rf_feature_df = pd.DataFrame({
    'feature': feature_names,
    'importance': rf_importance
}).sort_values('importance', ascending=False)

# Gradient Boosting feature importance
gb_importance = gradient_boosting.feature_importances_
gb_feature_df = pd.DataFrame({
    'feature': feature_names,
    'importance': gb_importance
}).sort_values('importance', ascending=False)

# Plot feature importance comparison
fig, axes = plt.subplots(1, 2, figsize=(16, 8))

# Top 15 features from Random Forest
top_rf_features = rf_feature_df.head(15)
axes[0].barh(range(len(top_rf_features)), top_rf_features['importance'], color='lightblue')
axes[0].set_yticks(range(len(top_rf_features)))
axes[0].set_yticklabels(top_rf_features['feature'])
axes[0].set_xlabel('Importance')
axes[0].set_title('Top 15 Features - Random Forest')
axes[0].invert_yaxis()

# Top 15 features from Gradient Boosting
top_gb_features = gb_feature_df.head(15)
axes[1].barh(range(len(top_gb_features)), top_gb_features['importance'], color='lightcoral')
axes[1].set_yticks(range(len(top_gb_features)))
axes[1].set_yticklabels(top_gb_features['feature'])
axes[1].set_xlabel('Importance')
axes[1].set_title('Top 15 Features - Gradient Boosting')
axes[1].invert_yaxis()


plt.tight_layout()
plt.show()
```

```

print("Top 10 Most Important Features (Random Forest):")
print_rf_feature_df.head(10).to_string(index=False))

print("\nTop 10 Most Important Features (Gradient Boosting):")
print_gb_feature_df.head(10).to_string(index=False))

```



Top 10 Most Important Features (Random Forest):

	feature	importance
	Destination Port	0.067340
	source_file	0.065975
	Init_Win_bytes_backward	0.055071
	Bwd Packets/s	0.029580
	Flow IAT Mean	0.029441
	Flow IAT Std	0.027969
	Packet Length Mean	0.027132
	Fwd IAT Std	0.025926
	Init_Win_bytes_forward	0.025565
	Fwd IAT Mean	0.024942

Top 10 Most Important Features (Gradient Boosting):

	feature	importance
	Idle Mean	0.240179
	Bwd Packet Length Std	0.109400
	act_data_pkt_fwd	0.080240
	Average Packet Size	0.069569
	Bwd Packet Length Mean	0.069367
	PSH Flag Count	0.066146
	Bwd Header Length	0.063291

Total Length of Fwd Packets	0.058635
Min Packet Length	0.026591
Max Packet Length	0.021086

6 6. MLflow Tracking

This block configures MLflow by setting the tracking URI and experiment name, then starts a run named “test run” to automatically log all trained models from the models list. It uses the appropriate MLflow logging function (mlf.sklearn.log_model or mlf.xgboost.log_model) to save each model as an artifact and registers it for version tracking.

```
[ ]: mlf.set_tracking_uri(uri="http://localhost:5001")
mlf.set_experiment("Cybersecurity project")
with mlf.start_run(run_name = "test run"):
    for mod in models:
        if mod["library"] == "sklearn":
            mlf.sklearn.log_model(
                sk_model=mod["model"],
                artifact_path=mod["name"],
                registered_model_name=f"reg_{mod['name']}")

        else:
            mlf.xgboost.log_model(
                xgb_model=mod["model"],
                artifact_path=mod["name"],
                registered_model_name=f"reg_{mod['name']}")
```

```
2025/12/09 21:27:35 INFO mlflow.tracking.fluent: Experiment with name
'Cybersecurity project' does not exist. Creating a new experiment.
2025/12/09 21:27:36 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/12/09 21:27:49 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.
Successfully registered model 'reg_log_reg'.
2025/12/09 21:27:51 INFO mlflow.store.model_registry.abstract_store: Waiting up
to 300 seconds for model version to finish creation. Model name: reg_log_reg,
version 1
Created version '1' of model 'reg_log_reg'.
2025/12/09 21:27:51 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/12/09 21:27:55 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.
Successfully registered model 'reg_log_reg_smote'.
2025/12/09 21:27:55 INFO mlflow.store.model_registry.abstract_store: Waiting up
```

```
to 300 seconds for model version to finish creation. Model name:  
reg_log_reg_smote, version 1  
Created version '1' of model 'reg_log_reg_smote'.  
2025/12/09 21:27:55 WARNING mlflow.models.model: `artifact_path` is deprecated.  
Please use `name` instead.  
2025/12/09 21:27:59 WARNING mlflow.models.model: Model logged without a  
signature and input example. Please set `input_example` parameter when logging  
the model to auto infer the model signature.  
Successfully registered model 'reg_knn'.  
2025/12/09 21:28:00 INFO mlflow.store.model_registry.abstract_store: Waiting up  
to 300 seconds for model version to finish creation. Model name: reg_knn,  
version 1  
Created version '1' of model 'reg_knn'.  
2025/12/09 21:28:00 WARNING mlflow.models.model: `artifact_path` is deprecated.  
Please use `name` instead.  
2025/12/09 21:28:04 WARNING mlflow.models.model: Model logged without a  
signature and input example. Please set `input_example` parameter when logging  
the model to auto infer the model signature.  
Successfully registered model 'reg_knn_smote'.  
2025/12/09 21:28:27 INFO mlflow.store.model_registry.abstract_store: Waiting up  
to 300 seconds for model version to finish creation. Model name: reg_knn_smote,  
version 1  
Created version '1' of model 'reg_knn_smote'.  
2025/12/09 21:28:27 WARNING mlflow.models.model: `artifact_path` is deprecated.  
Please use `name` instead.  
2025/12/09 21:28:31 WARNING mlflow.models.model: Model logged without a  
signature and input example. Please set `input_example` parameter when logging  
the model to auto infer the model signature.  
Successfully registered model 'reg_rand_forest'.  
2025/12/09 21:28:32 INFO mlflow.store.model_registry.abstract_store: Waiting up  
to 300 seconds for model version to finish creation. Model name:  
reg_rand_forest, version 1  
Created version '1' of model 'reg_rand_forest'.  
2025/12/09 21:28:32 WARNING mlflow.models.model: `artifact_path` is deprecated.  
Please use `name` instead.  
2025/12/09 21:28:36 WARNING mlflow.models.model: Model logged without a  
signature and input example. Please set `input_example` parameter when logging  
the model to auto infer the model signature.  
Successfully registered model 'reg_rand_forest_smote'.  
2025/12/09 21:28:41 INFO mlflow.store.model_registry.abstract_store: Waiting up  
to 300 seconds for model version to finish creation. Model name:  
reg_rand_forest_smote, version 1  
Created version '1' of model 'reg_rand_forest_smote'.  
2025/12/09 21:28:41 WARNING mlflow.models.model: `artifact_path` is deprecated.  
Please use `name` instead.  
2025/12/09 21:28:47 WARNING mlflow.models.model: Model logged without a  
signature and input example. Please set `input_example` parameter when logging  
the model to auto infer the model signature.
```

```
Successfully registered model 'reg_XGB'.
2025/12/09 21:28:49 INFO mlflow.store.model_registry.abstract_store: Waiting up
to 300 seconds for model version to finish creation. Model name: reg_XGB,
version 1
Created version '1' of model 'reg_XGB'.
2025/12/09 21:28:49 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/12/09 21:28:54 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.
Successfully registered model 'reg_XGB_smote'.
2025/12/09 21:28:54 INFO mlflow.store.model_registry.abstract_store: Waiting up
to 300 seconds for model version to finish creation. Model name: reg_XGB_smote,
version 1
Created version '1' of model 'reg_XGB_smote'.
2025/12/09 21:28:54 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/12/09 21:28:59 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.
Successfully registered model 'reg_hard_voting'.
2025/12/09 21:30:06 INFO mlflow.store.model_registry.abstract_store: Waiting up
to 300 seconds for model version to finish creation. Model name:
reg_hard_voting, version 1
Created version '1' of model 'reg_hard_voting'.
2025/12/09 21:30:06 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/12/09 21:30:25 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.
Successfully registered model 'reg_soft_voting'.
2025/12/09 21:31:08 INFO mlflow.store.model_registry.abstract_store: Waiting up
to 300 seconds for model version to finish creation. Model name:
reg_soft_voting, version 1
Created version '1' of model 'reg_soft_voting'.
2025/12/09 21:31:09 WARNING mlflow.models.model: `artifact_path` is deprecated.
Please use `name` instead.
2025/12/09 21:31:13 WARNING mlflow.models.model: Model logged without a
signature and input example. Please set `input_example` parameter when logging
the model to auto infer the model signature.
Successfully registered model 'reg_stacking'.
2025/12/09 21:32:19 INFO mlflow.store.model_registry.abstract_store: Waiting up
to 300 seconds for model version to finish creation. Model name: reg_stacking,
version 1
```

View run test run at: <http://localhost:5001/#/experiments/312917764447112075/runs/5acd235c9dde4a518023b614fe5b4f86>

View experiment at: <http://localhost:5001/#/experiments/312917764447112075>

Created version '1' of model 'reg_stacking'.